

Users, Files, and the Manual

File I/O

Prof. Seokin Hong

Kyungpook National University

Fall 2019

Objectives

■ Ideas and Skills

- The role and use of on-line documentation
- The Unix file interface: open, read, write, lseek, close
- Reading, creating, and writing files
- File descriptors
- Buffering: user level and kernel level
- Kernel mode, user mode, and the cost of system calls
- How Unix represents time, how to format Unix time
- Using the utmp file to find list of current users
- Detecting and reporting errors in system calls

■ System Calls and Functions

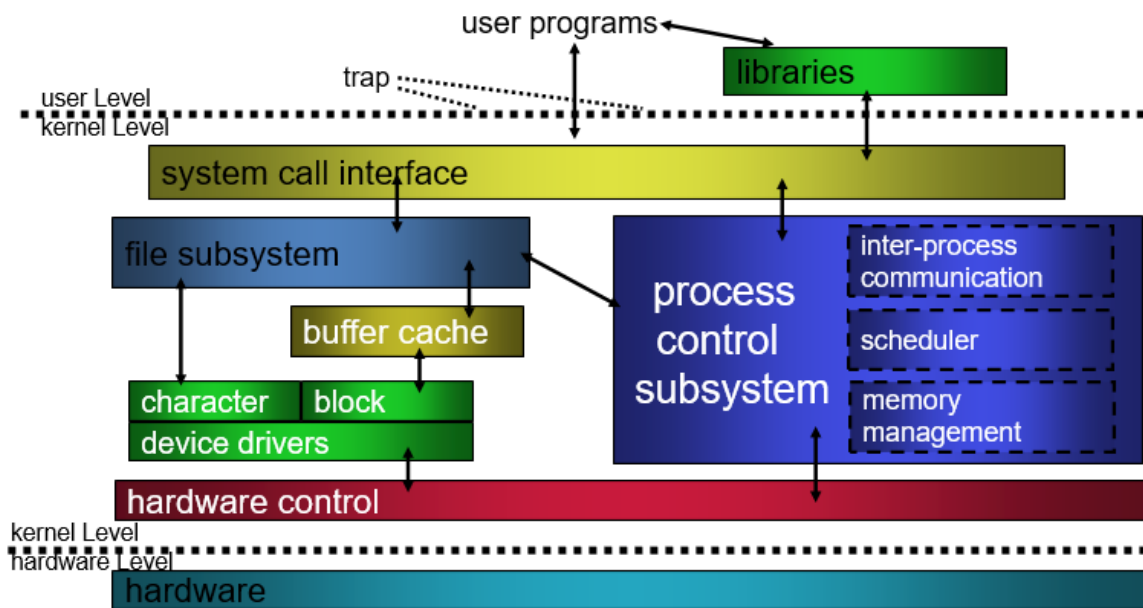
- open, read, write, create, lseek, close
- perror

■ UNIX/LINUX Commands

- man, who, cp, login

System Call

- **System calls** are an operating system's API
 - The set of functions that the operating system exposes to processes
- If you want to the OS to do something, you tell it via a **system call**



What are System Calls used for?

- Anything to do with
 - Accessing devices
 - Accessing files
 - Requesting memory
 - Setting/Changing access permissions
 - Communicating with other processes
 - Stopping/starting processes
 - Setting a timer

- You need a system call to
 - **Open a file**
 - Get data from the network
 - Kill a process

Three step to learn Unix/Linux system programming

- Looking at “real” programs : What does that do?
- Looking at the system calls : How does that work?
- Writing our own version : Can I try to do it?

Agenda

- 2.2 Asking about who
- 2.3 What Does who Do?
- 2.4 How Does who Do It?
- 2.5 Can I Write who?
- 2.6 Writing cp (read and write)
- 2.7 More Efficient File I/O: Buffering
- 2.8 Buffering and the Kernel
- 2.9 Reading and Writing a File
- 2.10 What to Do with System-Call Errors

Asking About **who**

- What does **who** do?
- How does **who** work?
- Can I write **who**?




FIGURE 2.1

Users, files, processes, devices, and kernel.

Commands are Programs

- Almost **all Unix commands** are **simply programs** written by a variety of people, usually in C.
- When you type **ls**, you are asking your **shell** to run the program named **ls**.
- Adding new commands to Unix is easy;
You write a new program and have the executable file stored in one of the standard system directories:
/bin, /usr/bin, /usr/local/bin.

Q1: What Does **who** Do?

- The **who** command displays information about users and processes on the local system

```
$ who
```

hecker1	ttyp1	Jul 21 19:51	(tide75.surfcity.com)
nlopez	ttyp2	Jul 21 18:11	(roam163-141.student.ivy.edu)
dgsulliv	ttyp3	Jul 21 14:18	(h004005a8bd64.ne.mediaone.net)
ackerman	ttyp4	Jul 15 22:40	(asd1-254.fas.state.edu)
wwchen	ttyp5	Jul 21 19:57	(circle.square.edu)
barbier	ttyp6	Jul 8 13:08	(labpc18.elsie.special.edu)
ramakris	ttyp7	Jul 13 08:51	(roam157-97.student.ivy.edu)
czhu	ttyp8	Jul 21 12:47	(spa.sailboat.edu)
bpsteven	ttyp9	Jul 21 18:26	(207.178.203.99)
molay	ttypa	Jul 21 20:00	(xyz73-200.harvard.edu)

\$

username

terminal name

login time

login host

one log-in session

What does **who** do?

Reading the Manual

- Every Unix system comes with manual for all commands
- The manual is on the disk and the command to read a page from the manual is **man**

```
$ man who
who(1)

NAME

    who - Identifies users currently logged in

...
SYNOPSIS

    who [-a]  | [-AbdhHlmMpqrstTu] [file]

    who am i

    who am I

    whoami
```

Q2: How Does **who** Do It?

- To learn more about Unix
 - Read the manual
 - Search the manual
 - Read the .h files
 - Follow SEE ALSO links

Read the Manual

- `$ man who`

DESCRIPTION

The `who` utility can list the user's name, terminal line, login time, elapsed time since activity occurred on the line, and the process-ID of the command interpreter (shell) for each current UNIX system user. It examines the `/var/adm/utmp` file to obtain its information. If file is given, that file (which must be in `utmp(4)` format) is examined. Usually, file will be `/var/adm/wtmp`, which contains a history of all the logins since the file was last created.

Read the .h files

- \$ more /usr/include/utmp.h

```
#define UTMP_FILE      "/var/adm/utmp"
#define WTMP_FILE      "/var/adm/wtmp"

#include <sys/types.h>  /* for pid_t, time_t */

/*
 * Structure of utmp and wtmp files.
 *
 * Assuming these numbers is unwise.
 */

#define ut_name ut_user      /* compatibility */
struct utmp {
    char    ut_user[32];      /* User login name */
    char    ut_id[14];        /* /etc/inittab id- IDENT_LEN in
                               * init */
    char    ut_line[32];      /* device name (console, lnxx) */
    short   ut_type;          /* type of entry */
    pid_t   ut_pid;           /* process id */
    struct exit_status {
        short    e_termination; /* Process termination status */
        short    e_exit;         /* Process exit status */
    } ut_exit;                /* The exit status of a process
                               * marked as DEAD_PROCESS.
                               */
    time_t   ut_time;         /* time entry was made */
    char     ut_host[64];      /* host name same as
                               * MAXHOSTNAMELEN */
};
/* Definitions for ut_type */
```

utmp.h (60%)

We now know how who works

- who works by:

```
open utmp  
[→ read record ←]  
[→ display record ←]  
close utmp←
```




FIGURE 2.2

Data flow in the `who` command.

utmp

Can I Write **who**?

- Two tasks we need to program
 - Read structs from a file
 - Display the information stored in a struct

We use open, read, and close

- Opening a file: **open()**
 - Opening a file is a **kernel service**;
 - **open() system call** is a request from your program to the kernel

```
fd = open(name, mode)
      ↑      ↙
char *  O_RDONLY
        O_WRONLY
        O_RDWR

returns -1 on error
OR an int on success
```




FIGURE 2.3

A file descriptor is a connection to a file.

More about open()

open		
PURPOSE	Creates a connection to a file	
INCLUDE	#include <fcntl.h>	
USAGE	int fd = open(char *name, int how)	
ARGS	name	name of file
	how	O_RDONLY, O_WRONLY, or O_RDWR
RETURNS	-1	on error
	int	on success

We use **open**, **read**, and **close**

■ Reading Data from a File: **read()**

```
fd = open(name, mode)
```

```
n = read(fd, array, numchars)
```

read		
PURPOSE	Transfer up to qty bytes from fd to buf	
INCLUDE	#include <unistd.h>	
USAGE	ssize_t numread = read(int fd, void *buf, size_t qty)	
ARGS	fd	source of data
	buf	destination for data
	qty	number of bytes to transfer
RETURNS	-1	on error
	numread	on success

We use `open`, `read`, and `close`

■ Closing a File: `close()`

```
fd = open(name, mode)
```

```
n = read(fd, array, numchars)
```

```
close(fd)
```

close	
PURPOSE	Closes a file
INCLUDE	#include <unistd.h>
USAGE	int result = close(int fd)
ARGS	fd file descriptor
RETURNS	-1 on error 0 on success

Why we need to close a file?

Writing who1.c

```
/* who1.c - a first version of the who program
 *
 *      open, read UTMP file, and show results
 */
#include      <stdio.h>
#include      <utmp.h>
#include      <fcntl.h>
#include      <unistd.h>
#include      <stdlib.h>
#define SHOWHOST      /* include remote machine on output */
void show_info( struct utmp* );
int main()
{
    struct utmp      current_record; /* read info into here      */
    int              utmpfd;          /* read from this descriptor */
    int              reclen = sizeof(current_record);

    if ( (utmpfd = open(UTMP_FILE, O_RDONLY)) == -1 ){
        perror( UTMP_FILE );      /* UTMP_FILE is in utmp.h  */
        exit(1);
    }
    while ( read(utmpfd, &current_record, reclen) == reclen )
        show_info(&current_record);
    close(utmpfd);
    return 0;                      /* went ok */
}
```

Displaying Log-In Records

```
/*
 *  show_info()
 *      displays contents of the utmp struct in human readable form
 *      *note* these sizes should not be hardwired
 */
void show_info( struct utmp* utbufp);
{
    printf("%-8.8s", utbufp->ut_name);        /* the logname */
    printf(" ");                             /* a space */
    printf("%-8.8s", utbufp->ut_line);        /* the tty */
    printf(" ");                             /* a space */
    printf("%10ld", utbufp->ut_time);         /* login time */
    printf(" ");                             /* a space */
#ifdef SHOWHOST
    printf("(%s)", utbufp->ut_host);          /* the host */
#endif
    printf("\n");                            /* newline */
}
```

Compile and run it:

```
$ cc who1.c -o who1
```

```
$ ./who1
```

```
      system b  952601411 ()
      run-leve  952601411 ()
                    952601416 ()
                    952601416 ()
                    952601417 ()
                    952601417 ()
                    952601419 ()
                    952601419 ()
                    952601423 ()
                    952601566 ()
LOGIN      console  952601566 ()
            ttypl   958240622 ()
shpyrko    ttypp2   964318862 (nas1-093.gas.swamp.org)
acotton    ttypp3   964319088 (math-guest04.williams.edu)
            ttypp4   964320298 ()
spradlin   ttypp5   963881486 (h002078c6adfb.ne.rusty.net)
dkoh       ttypp6   964314388 (128.103.223.110)
spradlin   ttypp7   964058662 (h002078c6adfb.ne.rusty.net)
king       ttypp8   964279969 (blade-runner.mit.edu)
berschba   ttypp9   964188340 (dudley.learned.edu)
rserved    ttyppa   963538145 (gigue.eas.ivy.edu)
dabel      ttyppb   964319455 (roam193-27.student.state.edu)
            ttyppc   964319645 ()
```

Let's compare our program with the system version

```
$ who
shpyrko      ttyp2      Jul 22 22:21      (nas1-093.gas.swamp.edu)
acotton      ttyp3      Jul 22 22:24      (math-guest04.williams.edu)
spradlin     ttyp5      Jul 17 20:51      (h002078c6adfb.ne.rusty.net)
dkoh         ttyp6      Jul 22 21:06      (128.103.223.110)
spradlin     ttyp7      Jul 19 22:04      (h002078c6adfb.ne.rusty.net)
king         ttyp8      Jul 22 11:32      (blade-runner.mit.edu)
berschba     ttyp9      Jul 21 10:05      (dudley.learned.edu)
rserved      ttya       Jul 13 21:29      (gigue.eas.ivy.edu)
dabel        ttypb      Jul 22 22:30      (roam193-27.student.state.edu)
rserved      ttyd       Jul 13 21:31      (gigue.eas.harvard.edu)
dkoh         ttye       Jul 22 16:46      (128.103.223.110)
molay        ttyq0      Jul 22 20:03      (xyz73-200.harvard.edu)
cweiner      ttyq8      Jul 21 16:40      (roam175-157.student.stats.edu)
$
```

■ What We Need to Do

- Suppress blank records
- Get the log-in times correct

Writing `who2.c`

■ Suppressing blank records

○ `/usr/include/utmp.h`

```
/*      Definitions for ut_type                                */
#define EMPTY          0
#define RUN_LVL        1
#define BOOT_TIME      2
#define OLD_TIME       3
#define NEW_TIME       4
#define INIT_PROCESS   5      /* Process spawned by "init" */
#define LOGIN_PROCESS  6      /* A "getty" process waiting for login */
#define USER_PROCESS    7      /* A user process */
#define DEAD_PROCESS   8
```

✗ represents the user logged into the system.

○ modification

```
show_info( struct utmp *utbufp )
{
    if ( utbufp->ut_type != USER_PROCESS )      /* users only ! */
        return;
    printf("%-8.8s", utbufp->ut_name);           /* the username */
}
```


Writing `who2.c` (cont.)

- Displaying Log-in Time in Human-Readable Form
 - How unix stores times: `time_t`
 - Unix stores times as the number of seconds since midnight, Jan 1, 1970, G.M.T. The `time_t` data type is an integer that stores a number of seconds.
 - Making a `time_t` readable: `ctime`

Writing who2.c (cont.)

■ Displaying Log-in Time in Human-Readable Form

```
$ man 3 ctime
```

```
CTIME(3)
```

```
Linux Programmer's Manual
```

```
CTIME(3)
```

```
NAME
```

```
asctime, ctime, gmtime, localtime, mktime - transform  
binary date and time to ASCII
```

```
SYNOPSIS
```

```
#include <time.h>
```

```
char *asctime(const struct tm *timeptr);
```

```
char *ctime(const time_t *timep);
```

```
...
```

The ctime() function converts the calendar time timep into a string of the form

```
"Wed Jun 30 21:49:08 1993\n"
```

Writing `who2.c` (cont.)

■ Putting it All together

```
/* who2.c - read /var/adm/utmp and list info therein
 *          - suppresses empty records
 *          - formats time nicely
 */
#include      <stdio.h>
#include      <unistd.h>
#include      <utmp.h>
#include      <fcntl.h>
#include      <time.h>
#include      <stdlib.h>
/* #define SHOWHOST */

void showtime(long);
void show_info(struct utmp *);

int main()
{
    struct utmp    utbuf;          /* read info into here */
    int            utmpfd;         /* read from this descriptor */
    ⌋ ✖ $ ls -l /usr/include | more
```

Can I Write [who](#)?

```
    if ( (utmpfd = open(UTMP_FILE, O_RDONLY)) == -1 ){
        perror(UTMP_FILE);
        exit(1);
    }

    while( read(utmpfd, &utbuf, sizeof(utbuf)) == sizeof(utbuf) )
        show_info( &utbuf );
    close(utmpfd);
    return 0;
}
/*
 *      show info()
 *
 *          displays the contents of the utmp struct
 *          in human readable form
 *          * displays nothing if record has no user name
 */
void show_info( struct utmp *utbufp )
{
    if ( utbufp->ut_type != USER_PROCESS )
        return;

    printf("%-8.8s", utbufp->ut_name);      /* the logname */
    printf(" ");                          /* a space */
    printf("%-8.8s", utbufp->ut_line);      /* the tty */
    printf(" ");                          /* a space */
    showtime( utbufp->ut_time );           /* display time */
#ifdef SHOWHOST
    if ( utbufp->ut_host[0] != '\0' )
        printf(" (%s)", utbufp->ut_host); /* the host */
#endif
    printf("\n");                          /* newline */
}
```

Can I Write `who`?

```
void showtime( long timeval )
/*
 *   displays time in a format fit for human consumption
 *   uses ctime to build a string then picks parts out of it
 *   Note: %12.12s prints a string 12 chars wide and LIMITS
 *   it to 12chars.
 */
{
    char    *cp;                /* to hold address of time    */
    cp = ctime(&timeval);       /* convert time to string    */
                                /* string looks like         */
                                /* Mon Feb  4 00:46:40 EST 1991 */
                                /* 0123456789012345.         */
    printf("%12.12s", cp+4 );    /* pick 12 chars from pos 4  */
}
```

Wed Jun 30 21:49:08 1993\n

partial string!

Testing who2.c

```
$ cc who2.c -o who2
```

```
$ ./who2
```

```
rlscott  tty2    Jul 23 01:07
acotton  tty3    Jul 22 22:24
spradlin tty5    Jul 17 20:51
spradlin tty7    Jul 19 22:04
king     tty8    Jul 22 11:32
berschba tty9    Jul 21 10:05
rserverd ttypa   Jul 13 21:29
rserverd ttyqd   Jul 13 21:31
molay    ttyq0   Jul 22 20:03
cweiner  ttyq8   Jul 21 16:40
mnabavi  ttyx2   Apr 10 23:11
```

```
$ who
```

```
rlscott      tty2      Jul 23 01:07
acotton      tty3      Jul 22 22:24
spradlin     tty5      Jul 17 20:51
spradlin     tty7      Jul 19 22:04
king         tty8      Jul 22 11:32
berschba     tty9      Jul 21 10:05
rserverd     ttypa     Jul 13 21:29
rserverd     ttyqd     Jul 13 21:31
molay        ttyq0     Jul 22 20:03
cweiner      ttyq8     Jul 21 16:40
mnabavi      ttyx2     Apr 10 23:11
```

```
$
```

Agenda

5 min break!

- 2.2 Asking about who
- 2.3 What Does who Do?
- 2.4 How Does who Do It?
- 2.5 Can I Write who?
- 2.6 Writing cp (read and write)
- 2.7 More Efficient File I/O: Buffering
- 2.8 Buffering and the Kernel
- 2.9 Reading and Writing a File
- 2.10 What to Do with System-Call Errors

Q1: What does **cp** do?

- **cp** makes a copy of a file

\$ **cp** source-file target-file

- If there is no target file, **cp** creates it.
- If there is a target file, **cp** replaces the contents of that file with the contents of the source file.

Q2: How Does cp Create and Write?

- Creating/Truncating a File:

```
fd = creat(name, 644);
```

creat		
PURPOSE	Create or zero a file	
INCLUDE	#include <fcntl.h>	
USAGE	int fd = creat(char *filename, mode_t mode)	
ARGS	filename:	the name of the file
	mode:	access permission
RETURNS	-1	on error
	fd	on success

Q2: How Does **cp** Create and Write?

■ Writing to a File


```
n = write(fd, buffer, num);
```

write		
PURPOSE	Send data from memory to a file	
INCLUDE	#include <unistd.h>	
USAGE	ssize_t result = write(int fd, void *buf, size_t amt)	
ARGS	fd	a file descriptor
	buf	an array
	amt	how many bytes to write
RETURNS	-1	on error
	num written	on success

Q3: Can I Write **cp**?

■ Program outline

```
open sourcefile for reading
open copyfile   for writing
+--> read from source to buffer -- eof? --+
|__ write from buffer to copy              |
                                           |
close sourcefile      <-----+
close copyfile
```



Why do we need to close a file?

FIGURE 2.4

Copying a file by reading and writing.

Can I Write `cp`

```
/** cp1.c
 *      version 1 of cp - uses read and write with tunable buffer size
 *
 *      usage: cp1 src dest
 */
#include      <stdio.h>
#include      <unistd.h>
#include      <fcntl.h>
#include      <stdlib.h>
#define BUFFERSIZE      4096
#define COPYMODE      0644

void oops(char *, char *);

main(int ac, char *av[])
{
    int      in_fd, out_fd, n_chars;
    char      buf[BUFFERSIZE];

    /* check args */
    if ( ac != 3 ){
        fprintf( stderr, "usage: %s source destination\n", *av);
        exit(1);
    }

    /* open files */

    if ( (in_fd=open(av[1], O_RDONLY)) == -1 )
        oops("Cannot open ", av[1]);

    if ( (out_fd=creat( av[2], COPYMODE)) == -1 )
        oops( "Cannot creat", av[2]);

    /* copy files */
```

```
while ( (n_chars = read(in_fd , buf, BUFFERSIZE)) > 0 )
    if ( write( out_fd, buf, n_chars ) != n_chars )
        oops("Write error to ", av[2]);
if ( n_chars == -1 )
    oops("Read error from ", av[1]);

/* close files */

if ( close(in_fd) == -1 || close(out_fd) == -1 )
    oops("Error closing files", "");
}

void oops(char *s1, char *s2)
{
    fprintf(stderr, "Error: %s ", s1);
    perror(s2);
    exit(1);
}
```

Compile and run it:

```
$ gcc cp1.c -o cp1
$ ./cp1 cp1 copy.of.cp1
$ ls -l cp1 copy.of.cp1
-rw-r--r--  1 bruce    bruce          37419 Jul 23 03:12 copy.of.cp1
-rwxrwxr-x  1 bruce    bruce          37419 Jul 23 03:08 cp1
$ cmp cp1 copy.of.cp1
$
```

How well does our program respond to errors?

```
$ cp1 xxx123 file1
Error: Cannot open xxx123: No such file or directory
$ cp1 cp1 /tmp
Error: Cannot creat /tmp: Is a directory
```

Agenda

- 2.2 Asking about who
- 2.3 What Does who Do?
- 2.4 How Does who Do It?
- 2.5 Can I Write who?
- 2.6 Writing cp (read and write)
- **2.7 More Efficient File I/O: Buffering**
- 2.8 Buffering and the Kernel
- 2.9 Reading and Writing a File
- 2.10 What to Do with System-Call Errors

Does the Size of the Buffer Matter?

■ Example:

Ex: Filesize = 2500 bytes

If buffer = 100 bytes then
copy requires read() and write() calls

If buffer = 1000 bytes then
copy requires read() and write() calls

■ System calls consume time:

buffer size	execution time in seconds
1	50.29
4	12.81
16	3.28
64	0.96
128	0.56
256	0.37
512	0.27
1024	0.22
2048	0.19
4096	0.18
8192	0.18
16384	0.18

cp1 copying a 5MB file

Why System Calls are Time Consuming?

- It runs various kernel functions, and it also requires a shift from USER MODE to KERNEL MODE and back;
- Not only does transferring data take time but mode change takes time
- **Thus, try to minimize system calls**




FIGURE 2.5
Control flow during system calls.

Does This Mean That who2.c is Inefficient?

■ Yes!

- Making one system call for each line of output makes as much sense as buying pizza by the slice or eggs one at a time
- who2.c use one system call for each utmp record




■ A better idea:

- Read in a bunch of records at once
- Then, process the ones in your local storage one by one

Adding Buffering to who2.c

- We make who2.c much more efficient by using buffering to reduce system calls



File buffering with utmplib

main calls a function in `utmplib.c` to get the next struct `utmp`.

Functions in `utmplib.c` read structs 16 at a time from the disk into an array.

The kernel is called only when all 16 are used up.

FIGURE 2.6

Buffering disk data in user space.

Revised Version: who3.c

```
/* who3.c - who with buffered reads
 *          - surpresses empty records
 *          - formats time nicely
 *          - buffers input (using utmplib)
 */
```

```
#include <stdio.h>
#include <sys/types.h>
#include <utmp.h>
#include <fcntl.h>
#include <time.h>
#include <stdlib.h>
#define SHOWHOST
```

```
void show_info(struct utmp *);
void showtime(time_t);
```

```
int main()
{
```

```
    struct utmp    *utbufp,          /* holds pointer to next rec    */
                   *utmp_next();      /* returns pointer to next    */
```

```
    if ( utmp_open( UTMP_FILE ) == -1 ){
        perror(UTMP_FILE);
        exit(1);
    }
```

```
    while ( ( utbufp = utmp_next() ) != ((struct utmp *) NULL) )
        show_info( utbufp );
    utmp_close( );
    return 0;
```

```
}
```


```
/*
 * show info()
```

```
...
```

```
struct utmp utbuf;
int utmpfd;

if( ( utmpfd = open(UTMP_FILE, O_RDONLY)) == -1 ) {
    perror( UTMP_FILE );
    exit(1);
}

while( read(utmpfd, &utbuf, sizeof(utbuf) ) == sizeof(utbuf) )
    show_info(&utbuf);
close(utmpfd);
```



■ The code for `utmplib.c`

```
/* utmplib.c - functions to buffer reads from utmp file
 *
 *      functions are
 *          utmp_open( filename )    - open file
 *          returns -1 on error
 *          utmp_next( )             - return pointer to next struct
 *          returns NULL on eof
 *          utmp_close()             - close file
 *
 *      reads NRECS per read and then doles them out from the buffer
 */
```

■ The code for utmplib.c (cont.)

```
#include <stdio.h>
#include <fcntl.h>
#include <sys/types.h>
#include <utmp.h>

#define NRECS 16
#define NULLUT ((struct utmp *)NULL)
#define UTSIZE (sizeof(struct utmp))

static char utmpbuf[NRECS * UTSIZE];          /* storage */
static int num_recs;                          /* num stored */
static int cur_rec;                          /* next to go */
static int fd_utm = -1;                      /* read from */

utmp_open( char *filename )
{
    fd_utm = open( filename, O_RDONLY );      /* open it */
    cur_rec = num_recs = 0;                  /* no recs yet */
    return fd_utm;                          /* report */
}

struct utmp *utmp_next()
{
    struct utmp *rec;
    if ( fd_utm == -1 )                      /* error ? */
        return NULLUT;
    if ( cur_rec==num_recs && utmp_reload()==0 ) /* any more ? */
        return NULLUT;
    /* get address of next record */
    rec = ( struct utmp *) &utmpbuf[cur_rec * UTSIZE];
    cur_rec++;
    return rec;
}
```

■ The code for `utmpplib.c` (cont.)

```
int utmp_reload()
/*
 *   read next bunch of records into buffer
 */
{
    int      amt_read;

                                /* read them in          */
    amt_read = read( fd_utmp , utmpbuf, NRECS * UTSIZE );

                                /* how many did we get? */
    num_recs = amt_read/UTSIZE;

                                /* reset pointer          */
    cur_rec  = 0;
    return num_recs;
}

utmp_close()
{
    if ( fd_utmp != -1 )          /* don't close if not    */
        close( fd_utmp );        /* open                  */
}
```

- Compile and Run

```
$ gcc who3.c umtplib.c -o who3
```

```
$ ./who3
```


Agenda

- 2.2 Asking about who
- 2.3 What Does who Do?
- 2.4 How Does who Do It?
- 2.5 Can I Write who?
- 2.6 Writing cp (read and write)
- 2.7 More Efficient File I/O: Buffering
- **2.8 Buffering and the Kernel**
- 2.9 Reading and Writing a File
- 2.10 What to Do with System-Call Errors

If Buffering Is So Smart, Why Doesn't the Kernel Do It?

■ It does!

- To save time, the kernel keeps copies of disk blocks in memory
- The `read()` copies data into a process buffer from a kernel buffer, and `write()` copies data from the process buffer to a kernel buffer



Consequences of Kernel Buffering

- Faster “disk” I/O
- Optimized disk writes
- Need to write buffers to disk before shutdown

FIGURE 2.7

Buffering disk data in the kernel.

Agenda

- 2.2 Asking about who
- 2.3 What Does who Do?
- 2.4 How Does who Do It?
- 2.5 Can I Write who?
- 2.6 Writing cp (read and write)
- 2.7 More Efficient File I/O: Buffering
- 2.8 Buffering and the Kernel
- **2.9 Reading and Writing a File**
- 2.10 What to Do with System-Call Errors

Logging Out: What It Does

- The system changes a record in the utmp file.
 - Specifically what changes?

- Experiment to see how it works:
 - 1. Log in to one machine.
 - 2. Use the who1 program we wrote to see the contents of utmp.
 - 3. Log out of one of your sessions.
 - 4. Repeat 1-3 to see what happened to the utmp record.

Logging Out: How It Works

- The program that removes your name from the log has to do the following:

- **1. Open the utmp file.**

```
fd = open(UTMP_FILE, O_RDWR);
```

- **2. Read the utmp file until it finds the record for your terminal.**

```
While( read(fd, rec, utmplen) == utmplen ) /* get next record */  
    If( strcmp(rec.ut_line, myline) == 0) /* what, my line? */  
        revise_entiry(); /* remove my name */
```

Logging Out: How It Works (cont.)

- The program that removes your name from the log has to do the following:
 - **3. Write a revised utmp record in its place.**
 - How do we write the revised record back to the file?
 - › If we just call write, our code updates the next record

Q: How can a program change the current read-write position in a file?

A: The lseek system call.

- **4. Close the utmp file.**

```
close(fd);
```

Moving the Current Position: `lseek()`

- Every open file has a current position
 - The current position belongs to the connection to the file, not to the file

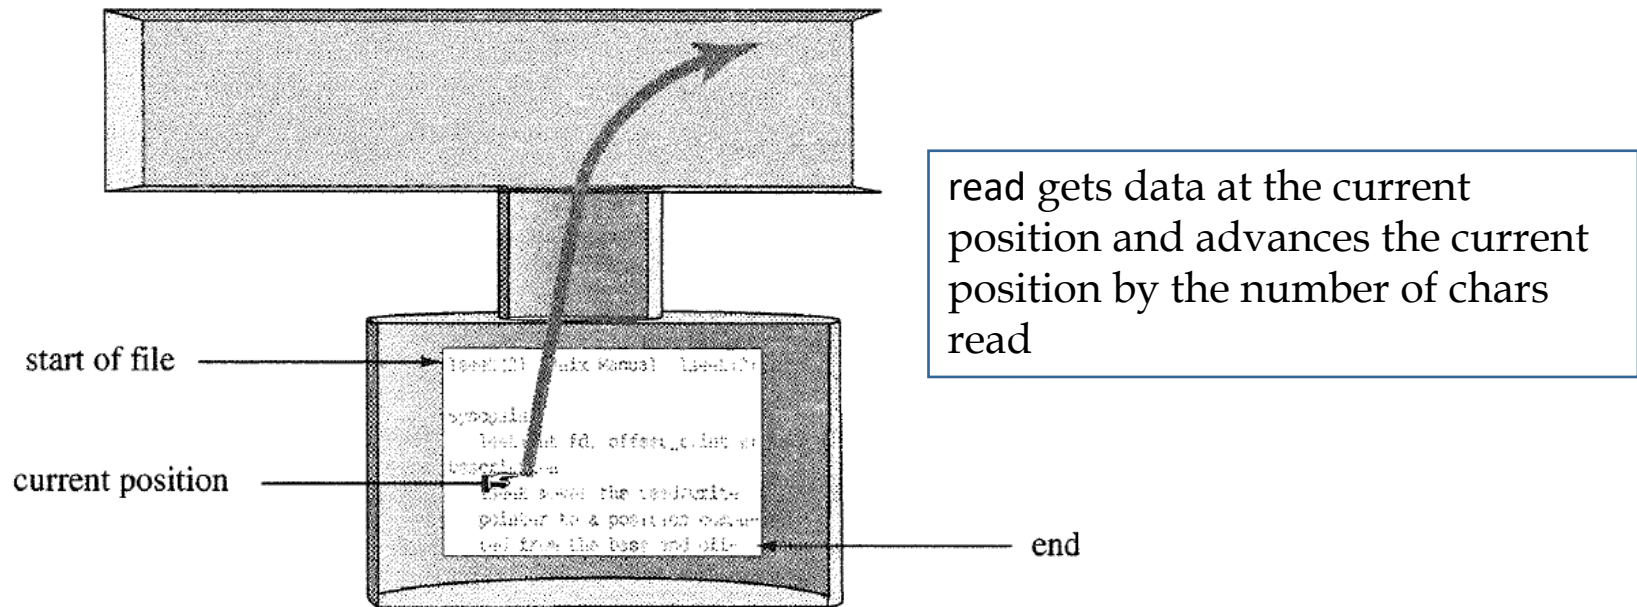



FIGURE 2.8

Every open file has a current position.

Moving the Current Position : lseek()

- offset: location in bytes
- whence: SEEK_SET, SEEK_CUR, SEEK_END



Moving the Current Position: `lseek()`

- The `lseek()` system call lets you change the current position of an open file:

```
lseek(fd, -(sizeof(struct utmp)), SEEK_CUR);  
lseek(fd, 10 * sizeof(struct utmp), SEEK_SET);
```

```
lseek(fd, 0, SEEK_END);  
write(fd, "hello", strlen("hello"));
```

```
lseek(fd, 0, SEEK_CUR)
```

Moving the Current Position: **lseek()**

lseek		
PURPOSE	Set file pointer to specified offset in file	
INCLUDE	#include <sys/types.h> #include <unistd.h>	
USAGE	off_t oldpos = lseek(int fd, off_t dist, int base)	
ARGS	fd:	file descriptor
	dist:	a distance in bytes
	base:	SEEK_SET => start of file
		SEEK_CUR => current position
		SEEK_END => end of file
RETURNS	-1	on error
	or	the previous position in the file

Code to Log Out from a Terminal

```
/*
 * logout_tty(char *line)
 *   marks a utmp record as logged out
 *   does not blank username or remote host
 *   returns -1 on error, 0 on success
 */
int logout_tty(char *line)
{
    int      fd;
    struct utmp rec;
    int      len = sizeof(struct utmp);
    int      retval = -1 ;                /* pessimism */

    if ( (fd = open(UTMP_FILE,O_RDWR)) == -1 ) /* open file */
        return -1;

    /* search and replace */
    while ( read(fd, &rec, len) == len)
        if ( strcmp(rec.ut_line, line, sizeof(rec.ut_line)) == 0)
        {
            rec.ut_type = DEAD_PROCESS;    /* set type */
            if ( time( &rec.ut_time ) != -1 ) /* and time */
                if ( lseek(fd, -len, SEEK_CUR) != -1 ) /* back up */
                    if ( write(fd, &rec, len) == len ) /* update */
                        retval = 0;                /* success! */
            break;
        }

    /* close the file */
    if ( close(fd) == -1 )
        retval = -1;
    return retval;
}
```

Agenda

- 2.2 Asking about who
- 2.3 What Does who Do?
- 2.4 How Does who Do It?
- 2.5 Can I Write who?
- 2.6 Writing cp (read and write)
- 2.7 More Efficient File I/O: Buffering
- 2.8 Buffering and the Kernel
- 2.9 Reading and Writing a File
- 2.10 What to Do with System-Call Errors

What to Do with System-Call Errors

- When open cannot open a file, it returns -1
- When read cannot read data, it returns -1
- When lseek cannot seek, it returns -1
- System calls return -1 when something goes wrong
- Your programs should test the return value of every system call they make and take intelligent action when errors occur

How to Identify What Went Wrong: `errno`

- The kernel tells your program the cause of the error by storing an error code in a **global variable** called `errno`
- The manpage for `errno(3)` and the file `<errno.h>` include the error-code symbols and numeric codes

```
$ man 3 errno
```

```
#define EPERM          1      /* Operation not permitted */
#define ENOENT         2      /* No such file or directory */
#define ESRCH          3      /* No such process */
#define EINTR          4      /* Interrupted system call */
#define EIO            5      /* I/O error */
```

Different Responses to Different Errors

```
#include <errno.h>

extern int errno;

int sample()
{
    int fd;
    fd = open("file", O_RDONLY);
    if ( fd == -1 )
    {
        printf("Cannot open file: ");
        if ( errno == ENOENT )
            printf("There is no such file.");
        else if ( errno == EINTR )
            printf("Interrupted while opening file.");

        else if ( errno == EACCESS )
            printf("You do not have permission to open file.");
        ...
    }
}
```

Reporting Errors: `perror(3)`

- If you want to print a message describing the error, you could test the value of `errno` and print different messages for different values.
- Print a system error message

```
int sample()
{
    int fd;
    fd = open("file", O_RDONLY);
    if ( fd == -1 )
    {
        perror("Cannot open file");
        return;
    }
    ...
}
```

Objectives

■ Ideas and Skills

- The role and use of on-line documentation ([man](#))
- The Unix file interface: [open](#), [read](#), [write](#), [lseek](#), [close](#)
- [Reading, creating, and writing files](#)
- [File descriptors](#)
- [Buffering](#): user level and kernel level
- Kernel mode, user mode, and the [cost of system calls](#)
- How Unix represents time, [how to format Unix time](#)
- Using the [utmp](#) file to find list of current users
- [Detecting and reporting errors](#) in system calls

■ System Calls and Functions

- [open](#), [read](#), [write](#), [create](#), [lseek](#), [close](#)
- [perror](#)

■ UNIX/LINUX Commands

- [man](#), [who](#), [cp](#), [login](#)

Next Time

- Chapter 3 : Directories and File Properties

Homework#1

- Complete who3 and cp1 programs
 - who3: who3.c
 - cp1 : cp1.c
- Compressing the source code of the programs to LMS
 - Compressed file name : student id_hw1
- We will just check if the programs are working
- Due: 11:59:59 pm, September 11, 2019

Appendix: **File descriptor**

- All system calls use file descriptor to refer to open files
 - regular files, pipes, FIFO, sockets, terminals, devices
- Each process has its own list of open fd
- Three standard fd opened by default
 - inherited from the shell

FD	Purpose	POSIX name	stdio stream
0	standard input	STDIN_FILENO	stdin
1	standard output	STDOUT_FILENO	stdout
2	standard error	STDERR_FILENO	stderr

Appendix: File descriptor

