

# Directories and File Properties

Prof. Seokin Hong

Kyungpook National University

Fall 2019

---

# Objectives

## ■ Ideas and Skills

- A directory is a list of files
- How to read a directory
- Types of files and how to determine the type of a file
- Properties of files and how to determine properties of a file
- Bit sets and bit masks
- User and group ID numbers

## ■ System Calls and Functions

- opendir, readdir, closedir, seekdir
- stat
- chmod, chown, utime
- rename

## ■ Commands

- ls

---


# Agenda

- What Does **ls** Do?
- Brief Review of the File System Tree
- How Dos Is Work?
- Can I Write Is?
- Writing Is -l
- Three Special Bits
- Setting and Modifying the Properties of a File

---

# Files and Directories (Review)

- **Single hierarchical directory structure**



- **File types**

- Regular(plain), device, pipe, socket, dir, sym link

- **Directory**

- Special file that has the table of filenames and the reference to those files

---

## ls

- lists names of files and reports file attributes
- Example:

```
$ ls
Makefile      docs      ls2.c      s.tar      statdemo.c  tail1.c
chap03       ls1.c     old_src    stat1.c    tail1
$
```

```
$ ls -l
total 108
-rw-rw-r--  2 bruce  users      345 Jul 29 11:05 Makefile
-rw-rw-r--  1 bruce  users    27521 Aug  1 12:14 chap03
drwxrwxr-x  2 bruce  users     1024 Aug  1 12:15 docs
Type&permission  links  owner      group      size    modified-date/time  name
```

# ls

- Listing other directories and reporting on other files

Asking ls about Other Directories and Their Files	
Example	Action
ls /tmp	list names of files in /tmp directory
ls -l docs	show attributes of files in docs directory
ls -l ../Makefile	show attributes of ../Makefile
ls *.c	list names of files matching pattern *.c

---

# Popular Command-Line Options

## ■ Options:

Command	Action
<code>ls -a</code>	shows <u>“.”-files</u>
<code>ls -lu</code>	shows last-read time
<code>ls -s</code>	shows size in blocks
<code>ls -t</code>	sorts in time order
<code>ls -F</code>	shows file types

`ls -al`

## ■ A remark on Dot-Files (hidden file)

- **ls** does not list the name of a file if the first character of the filename is a dot.
- Some programs use dot filenames in a user's home directory to store user preferences.

```
root@DESKTOP-K4MA2V5:~# ls -l
```

합계 24

```
drwxrwxrwx 0 root root 512 2월 14 10:24 .
-rw-rw-rw- 1 root root 27 2월 28 14:39 cat.test
-rw-rw-rw- 1 root root 9525 2월 27 16:46 etc.listing
-rwxrwxrwx 1 root root 8600 3월 12 19:06 hello
-rw-rw-rw- 1 root root 61 3월 12 19:06 hello.c
-rw-rw-rw- 1 root root 33 2월 21 15:05 test.c
-rw-rw-rw- 1 root root 0 2월 12 15:02 userlist
-rw-rw-rw- 1 root root 13 3월 8 10:09 vitest.txt
```

Current directory

```
root@DESKTOP-K4MA2V5:~# ls -al
```

합계 32

```
drwx----- 0 root root 512 3월 12 19:06 .
drwxr-xr-x 0 root root 512 8월 25 2017 ..
-rw----- 1 root root 96 2월 28 10:45 .bash_history
-rw-r--r-- 1 root root 3106 10월 23 2015 .bashrc
drwxrwxrwx 0 root root 512 3월 12 18:59 .nano
-rw-r--r-- 1 root root 148 8월 18 2015 .profile
drwxr-xr-x 0 root root 512 3월 12 18:58 .vim
-rw----- 1 root root 3370 3월 12 18:59 .viminfo
drwxrwxrwx 0 root root 512 2월 14 10:24 .
-rw-rw-rw- 1 root root 27 2월 28 14:39 cat.test
-rw-rw-rw- 1 root root 9525 2월 27 16:46 etc.listing
-rwxrwxrwx 1 root root 8600 3월 12 19:06 hello
-rw-rw-rw- 1 root root 61 3월 12 19:06 hello.c
-rw-rw-rw- 1 root root 33 2월 21 15:05 test.c
-rw-rw-rw- 1 root root 0 2월 12 15:02 userlist
-rw-rw-rw- 1 root root 13 3월 8 10:09 vitest.txt
```

Parent directory

- Dot-file (hidden file)
- Dot-files at home directory are typically used for user preferences of programs

```
root@DESKTOP-K4MA2V5:~#
```



---

## So, what does **ls** do?

- **ls** does two things

1. Lists the contents of directories
2. Displays information about files

- We need to learn:

1. How to list the contents of a directory
2. How to obtain and display properties of a file
3. How to determine if a name refers to a file or a directory

---

# Agenda

- What Does **ls** Do?
- Brief Review of the File System Tree
- How **Dos ls** Work?
- Can I Write **ls**?
- Writing **ls -l**
- Three Special Bits
- Setting and Modifying the Properties of a File

---

# File System Tree

- The **disk** is organized as a **tree of directories**, each of which contains files or directories.
- The commands **cd**, **pwd**, **ls** allow us to explore a file system




FIGURE 3.1

A tree of directories.

---

## How Dos **ls** Work?

### ■ Outline :

```
        open directory
+--> read entry          -end of dir?--+
|__ display file info    |
    close directory      <-----+

```

### ■ It looks just like the logic for who!

### ■ Difference?

- The main difference is that the `who` opens and reads *from a file* (*utmp*, *wtmp*)
- **ls** opens and reads its data *from a directory*.

---

## What is a Directory?

- A directory is a kind of file that contains a list of names of files and directories.
- Unlike a regular file, a directory never empty
  - Every directory contains two specific items: **.** **and** **..**
  - **dot(.)** is the name of the current directory,
  - **dotdot(..)** is the name of the directory one level up.

---

## Do open, read, and close work for directories?

- Answer 1: on old versions of Unix, that was the only way
  - On some versions of Unix, you still can, but not for all directories
  
- Answer 2: It is a bad idea to use open, read and close to list contents of directory. Why?
  - Unix allows various disk formats to appear as part of a single tree.
    - It supports Mac HFS, FAT, FAT32, lots of Unix flavors;
  - Thus, using **read** to process each type would require knowing the format of the records for each type of directory

---

## How do I read a Directory? (I)

- We read the entires by calling readdir()
- Each readdir() call returns a pointer to the next record, a variable of type struct dirent

```
struct dirent *readdir (  
    DIR *dir_pointer );
```

# How do I read a Directory? (II)

```
#include <dirent.h>
```

```
opendir(char *)  
  creates a connection,  
  returns a DIR *
```

```
readdir(DIR *)  
  reads next record,  
  returns a pointer  
  to a struct dirent
```

```
closedir(DIR *)  
  closes a connection
```




FIGURE 3.2

Reading entries from a directory.



---

# How do I read a Directory? (III)

File Formats

dirent(4)

## NAME

dirent - file system independent directory entry

## SYNOPSIS

```
#include <dirent.h>
```

## DESCRIPTION

Different file system types may have different directory entries. The `dirent` structure defines a file system independent directory entry, which contains information common to directory entries in different file system types. A set of these structures is returned by the [`getdents\(2\)`](#) system call.

The `dirent` structure is defined:

```
struct dirent {
    ino_t      d_ino;
    off_t      d_off;
    unsigned short d_reclen;
    char       d_name[1];
};
```

---

# Agenda

- What Does **ls** Do?
- Brief Review of the File System Tree
- How **Dos ls** Work?
- Can I Write **ls**?
- Writing **ls -l**
- Three Special Bits
- Setting and Modifying the Properties of a File

---

## Writing ls1.c

- Logic for listing a directory:

```
main()
    opendir
    while ( readdir )
        print d_name
    closedir
```

```

/** ls1.c
**  purpose - list contents of directory or directories
**  action - if no args, use . else list files in args
**/
#include      <stdio.h>
#include      <sys/types.h>
#include      <dirent.h>

void do_ls(char []);

main(int ac, char *av[])
{
    if ( ac == 1 )
        do_ls( "." );
    else
        while ( --ac ){
            printf("%s:\n", *++av );
            do_ls( *av );
        }
}

void do_ls( char dirname[] )
/*
 *   list files in directory called dirname
 */
{
    DIR          *dir_ptr;           /* the directory */
    struct dirent *direntp;          /* each entry   */

    if ( ( dir_ptr = opendir( dirname ) ) == NULL )
        fprintf(stderr, "ls1: cannot open %s\n", dirname);
    else
    {
        while ( ( direntp = readdir( dir_ptr ) ) != NULL )
            printf("%s\n", direntp->d_name );
        closedir(dir_ptr);
    }
}

```

---

## ■ Compile and run it:

```
$ cc -o ls1 ls1.c
```

```
$ ls1
```

```
.  
..  
s.tar  
tail1  
Makefile  
ls1.c  
ls2.c  
chap03  
old_src  
docs  
ls1  
stat1.c  
statdemo.c  
tail1.c
```

```
$ ls
```

```
Makefile      docs          ls1.c         old_src       stat1.c       tail1  
chap03        ls1           ls2.c         s.tar         statdemo.c    tail1.c  
$
```

✖

```
$ ./ls1
```

```
$ ./ls1 . /tmp /usr
```

---

# Agenda

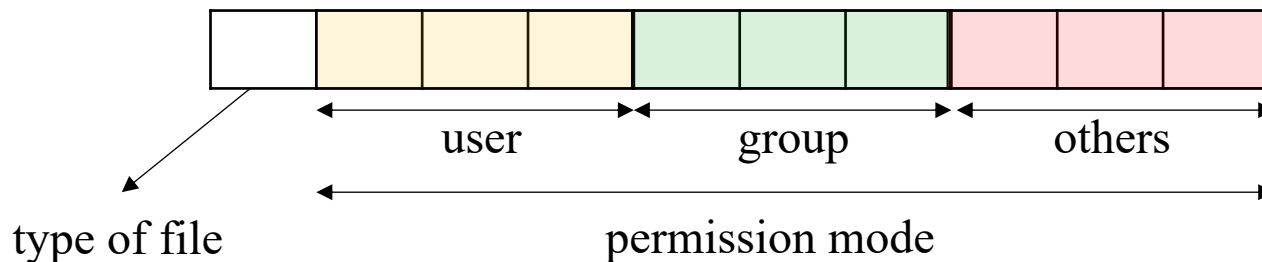
- What Does **ls** Do?
- Brief Review of the File System Tree
- How **Dos ls** Work?
- Can I Write **ls**?
- Writing **ls -l**
- Three Special Bits
- Setting and Modifying the Properties of a File

# What Does “ls -l” Do?

- **ls** does two different types of things
  - lists directories and files
  - display information about directories and files

```
$ ls -l
total 108
-rw-rw-r-- 2 bruce  users      345 Jul 29 11:05 Makefile
-rw-rw-r-- 1 bruce  users    27521 Aug  1 12:14 chap03
drwxrwxr-x 2 bruce  users     1024 Aug  1 12:15 docs
```


type and permission    links   owner        group                    size   last-modified time   name



- : regular file  
d : directory

# How Does “ls -l” work? (I)

- How can we get information (status/properties) about a file?
  - **stat** system call is used to retrieve file status
- How does **stat** work:
  - The process defines a buffer of type **struct state**
  - And then asks the kernel to copy file information from the disk to the buffer





---

## How Does “ls -l” work? (II)

### ■ struct stat

- Defined in /usr/include/sys/stat.h

```
struct stat {
    dev_t      st_dev;          /* ID of device containing file */
    ino_t      st_ino;          /* inode number */
    mode_t     st_mode;         /* protection */
    nlink_t    st_nlink;        /* number of hard links */
    uid_t      st_uid;          /* user ID of owner */
    gid_t      st_gid;          /* group ID of owner */
    dev_t      st_rdev;         /* device ID (if special file) */
    off_t      st_size;         /* total size, in bytes */
    blksize_t  st_blksize;      /* blocksize for filesystem I/O */
    blkcnt_t   st_blocks;       /* number of 512B blocks allocated */

    /* Since Linux 2.6, the kernel supports nanosecond
       precision for the following timestamp fields.
       For the details before Linux 2.6, see NOTES. */

    struct timespec st_atim;     /* time of last access */
    struct timespec st_mtim;     /* time of last modification */
    struct timespec st_ctim;     /* time of last status change */
}
```

---

## How Does “ls -l” work? (III)

stat		
PUPOSE	Obtain information about a file	
INCLUDE	#include <sys/stat.h>	
USAGE	int result = stat(char *fname, struct stat *bufp)	
AGRS	fname	name of file
	bufp	pointer to buffer
RETURNS	-1	if error
	0	if success

---

# Writing fileinfo.c

- Use stat to get file info for that name
- Display the items in the struct

```
/* fileinfo.c - use stat() to obtain and print file properties
 *             - some members are just numbers...
 */
#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>

void show_stat_info(char *, struct stat *);

int main(int ac, char *av[])
{
    struct stat info;          /* buffer for file info */

    if (ac>1)
        if( stat(av[1], &info) != -1 ){
            show_stat_info( av[1], &info );
            return 0;
        }
        else
            perror(av[1]); /* report stat() errors */
    return 1;
}
```

---

## Writing fileinfo.c

```
void show_stat_info(char *fname, struct stat *buf)
/*
 * displays some info from stat in a name=value format
 */
{
    printf("    mode: %o\n", buf->st_mode);          /* type + mode */
    printf("    links: %d\n", buf->st_nlink);        /* # links      */
    printf("    user: %d\n", buf->st_uid);           /* user id      */
    printf("    group: %d\n", buf->st_gid);          /* group id     */
    printf("    size: %d\n", buf->st_size);          /* file size    */
    printf("modtime: %d\n", buf->st_mtime);         /* modified     */
    printf("    name: %s\n", fname );              /* filename     */
}
```

---

## Writing fileinfo.c

- Compile and run it :

```
$ gcc -o fileinfo fileinfo.c
```

```
$ ./fileinfo fileinfo.c
```

```
mode: 100664
```

```
links: 1
```

```
user: 500
```

```
group: 120
```

```
size: 1106
```

```
modtime: 965158604
```

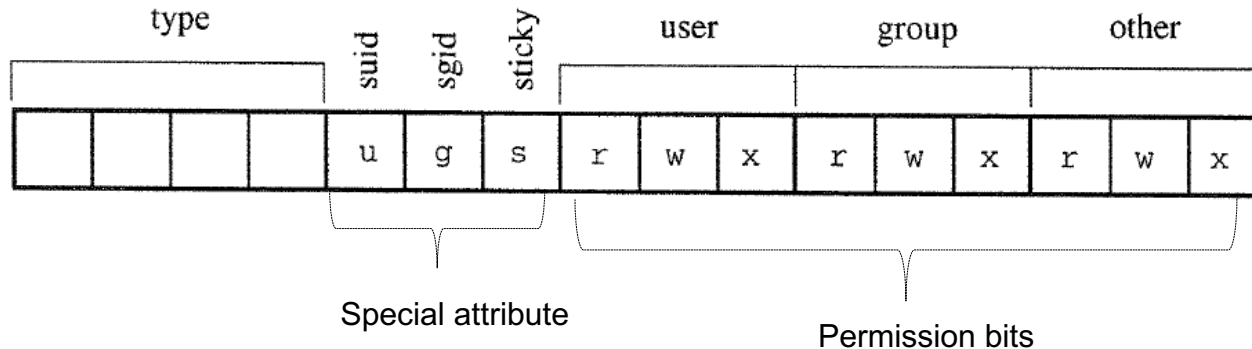
```
name: fileinfo.c
```

```
$ ls -l fileinfo.c
```

```
-rw-rw-r--  1 bruce  users          1106 Aug  1 15:36 fileinfo.c
```

# Converting file mode to a string

- File type and permission bits are stored in the **st\_mode** member



- **Type:** file types

- 4 bits means 16 possible patterns.
- Each pattern can correspond to a file type.

- **Permission bits :**

- Access permission of user, group, others for the file
- 1 indicates the permission is granted
- 0 indicates the permission is denied

```
$ ./fileinfo fileinfo.c
mode: 100664
```

# How to read subfields: Masking

- How do we examine a bit or sub-field?

- ex) 100664 (base 8) → -rw-rw-r--

- Use “bitwise AND (&)” to MASK

○ Ex)	1	0	0	0	0	0	0	1	1	0	1	1	0	1	0	0	value
&	0	0	0	0	0	0	0	1	0	1	1	0	0	0	0	0	mask
=	0	0	0	0	0	0	0	1	0	1	1	0	0	0	0	0	result

FIGURE 3.6

Applying a bitmask.

# Using Masking to decode permission bits

- 100664 (base 8) → -rw-rw-r--

```
/*
 * This function takes a mode value and a char array
 * and puts into the char array the file type and
 * nine letters that correspond to the bits of the
 * mode.
 * NOTE: It does not code setuid, setgid, and sticky
 * codes
 */
void mode_to_letters( int mode, char str[] )
{
    strcpy( str, "-----" );          /* default=no perms */
    if ( S_ISDIR(mode) ) str[0] = 'd';  /* directory? */
    if ( S_ISCHR(mode) ) str[0] = 'c';  /* char devices */
    if ( S_ISBLK(mode) ) str[0] = 'b';  /* block device */

    if ( mode & S_IRUSR ) str[1] = 'r';  /* 3 bits for user */
    if ( mode & S_IWUSR ) str[2] = 'w';
    if ( mode & S_IXUSR ) str[3] = 'x';

    if ( mode & S_IRGRP ) str[4] = 'r';  /* 3 bits for group */
    if ( mode & S_IWGRP ) str[5] = 'w';
    if ( mode & S_IXGRP ) str[6] = 'x';

    if ( mode & S_IROTH ) str[7] = 'r';  /* 3 bits for other */
    if ( mode & S_IWOTH ) str[8] = 'w';
    if ( mode & S_IXOTH ) str[9] = 'x';
}
```

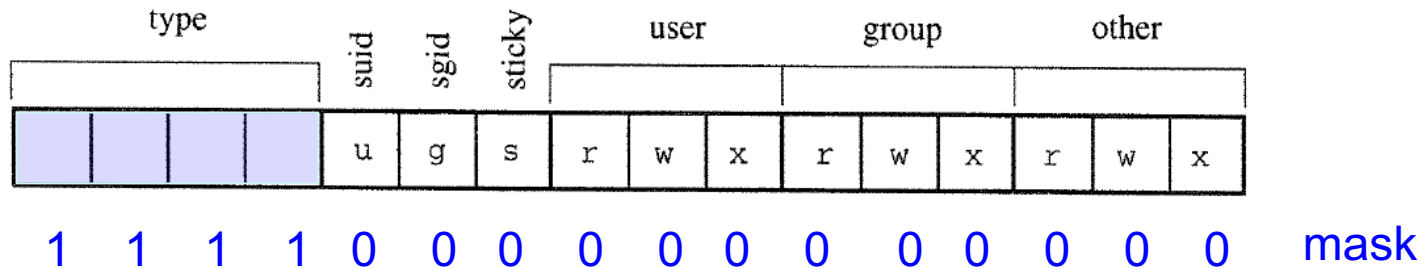
S_IRWXU	00700	owner has read, write, and execute permission
S_IRUSR	00400	owner has read permission
S_IWUSR	00200	owner has write permission
S_IXUSR	00100	owner has execute permission
S_IRWXG	00070	group has read, write, and execute permission
S_IRGRP	00040	group has read permission
S_IWGRP	00020	group has write permission
S_IXGRP	00010	group has execute permission

masks defined in <sys/stat.h>



# Using Masking to decode file types

- Mask and file types defined in <sys/stat.h>



```
#define S_IFMT      0170000    /* type of file */
#define S_IFREG     0100000    /* regular */
#define S_IFDIR     0040000    /* directory */
#define S_IFBLK     0060000    /* block special */
#define S_IFCHR     0020000    /* character special */
#define S_IFIFO     0010000    /* fifo */
#define S_IFLNK     0120000    /* symbolic link */
#define S_IFSOCK    0140000    /* socket */
```

File types

※ octal

```
if ( (info.st_mode & 0170000) == 0040000 )
    printf("this is a directory.");
```

---

# Using Masking to decode file types

```
drwxr-xr-x 2 root root    0 Jan  1  1970 home
```

## ■ File types

- **Regular** : regular file, marked with **-**
- **Directory** : directory. marked with **d**
- **Symbolic link** : a reference to another file. marked with **l**
- **Socket** : file used for inter-process communication that enable packetized-communication between two processes. communication can extend beyond localhost. marked with an **s**
- **Block special** : interface that allows an application to interact with a hardware devices. It provides buffered access to the hardware. marked with **b**
- **Character special** : interface that allows an application to interact with a hardware devices. It provides un-buffered, direct access to the hardware. marked with **c**
- **FIFO (named pipe)** : file used for inter-process communication within a host. marked with **p**

---

## Using **Macros** to decode file types

- **Macros** defined in `<sys/stat.h>`

```
/*
 *      File type macros
 */

#define S_ISFIFO(m)      (((m)&(0170000)) == (0010000))
#define S_ISDIR(m)       (((m)&(0170000)) == (0040000))
#define S_ISCHR(m)       (((m)&(0170000)) == (0020000))
#define S_ISBLK(m)       (((m)&(0170000)) == (0060000))
#define S_ISREG(m)       (((m)&(0170000)) == (0100000))

if ( S_ISDIR(info.st_mode) )
    printf("this is a directory.");
```

---

## Converting User ID to Strings

```
$ ./fileinfo fileinfo.c
  mode: 100664
  links: 1
  user: 500
  group: 120
  size: 1106
modtime: 965158604
  name: fileinfo.c
$ ls -l fileinfo.c
-rw-rw-r-- 1 bruce  users      1106 Aug  1 15:36 fileinfo.c
```

---

# Converting User ID to Strings

- Library function `getpwuid()` provides access to the complete list of users
  - Defined in `/usr/include/pwd.h`

```
struct passwd *getpwuid(uid_t uid);
```

- Example

```
char *uid_to_name( uid_t uid )
{
    return getpwuid(uid)->pw_name ;
}
```

- struct passwd

```
/* The passwd structure. */
struct passwd
{
    char *pw_name;           /* Username. */
    char *pw_passwd;         /* Password. */
    __uid_t pw_uid;          /* User ID. */
    __gid_t pw_gid;          /* Group ID. */
    char *pw_gecos;          /* Real name. */
    char *pw_dir;            /* Home directory. */
    char *pw_shell;          /* Shell program. */
};
```

---

## Converting Group ID to Strings

```
$ ./fileinfo fileinfo.c
```

```
mode: 100664
```

```
links: 1
```

```
user: 500
```

```
group: 120
```

```
size: 1106
```

```
modtime: 965158604
```

```
name: fileinfo.c
```

```
$ ls -l fileinfo.c
```

```
-rw-rw-r--  1 bruce  users      1106 Aug  1 15:36 fileinfo.c
```

---

# Converting Group ID to Strings

- `getgrgid()` provides access to the list of groups

- Defined in `/usr/include/grp.h`

```
struct group *getgrgid(gid_t gid);
```

- Example

```
char *gid_to_name( gid_t gid )  
{  
    return getgrgid(gid)->gr_name ;  
}
```

- `struct group`

```
struct group {  
    char    *gr_name;           /* group name */  
    char    *gr_passwd;        /* group password */  
    gid_t   gr_gid;            /* group ID */  
    char    **gr_mem;          /* group members */  
};
```

## Putting It All Together: ls2.c

```
$ cc -o ls1 ls1.c
```

```
$ ls1
```

```
.  
..  
s.tar  
tail1  
Makefile  
ls1.c  
ls2.c  
chap03  
old_src  
docs  
ls1  
stat1.c  
statdemo.c  
tail1.c
```

```
$ cc -o fileinfo fileinfo.c
```

```
$ ./fileinfo fileinfo.c
```

```
mode: 100664  
links: 1  
user: 500  
group: 120  
size: 1106  
modtime: 965158604  
name: fileinfo.c
```

```
$ ls2
```

drwxrwxr-x	4	bruce	bruce	1024	Aug	2	18:18	.
drwxrwxr-x	5	bruce	bruce	1024	Aug	2	18:14	..
-rw-rw-r--	1	bruce	users	30720	Aug	1	12:05	s.tar
-rwxrwxr-x	1	bruce	users	37351	Aug	1	12:13	tail1
-rw-rw-r--	2	bruce	users	345	Jul	29	11:05	Makefile
-rw-r--r--	1	bruce	users	723	Aug	1	14:26	ls1.c
-rw-r--r--	1	bruce	users	3045	Feb	15	03:51	ls2.c
-rw-rw-r--	1	bruce	users	27521	Aug	1	12:14	chap03
drwxrwxr-x	2	bruce	users	1024	Aug	1	12:14	old_src
drwxrwxr-x	2	bruce	users	1024	Aug	1	12:15	docs
-rwxrwxr-x	1	bruce	bruce	37048	Aug	1	14:26	ls1



---

```
/* ls2.c
 *      purpose  list contents of directory or directories
 *      action   if no args, use .  else list files in args
 *      note     uses stat and pwd.h and grp.h
 *      BUG: try ls2 /tmp
 */
#include      <stdio.h>
#include      <sys/types.h>
#include      <dirent.h>
#include      <sys/stat.h>
#include      <string.h>

void do_ls(char[]);
void dostat(char *);
void show_file_info( char *, struct stat *);
void mode_to_letters( int , char [] );
char *uid_to_name( uid_t );
char *gid_to_name( gid_t );

main(int ac, char *av[])
{
    if ( ac == 1 )
        do_ls( "." );
    else
        while ( --ac ){
            printf("%s:\n", *++av );
            do_ls( *av );
        }
}
```

---

```
void do_ls( char dirname[] )
/*
 *    list files in directory called dirname
 */
{
    DIR          *dir_ptr;          /* the directory */
    struct dirent *direntp;          /* each entry    */
    if ( ( dir_ptr = opendir( dirname ) ) == NULL )
        fprintf(stderr, "ls1: cannot open %s\n", dirname);
    else
    {
        while ( ( direntp = readdir( dir_ptr ) ) != NULL )
            dostat( direntp->d_name );
        closedir(dir_ptr);
    }
}

void dostat( char *filename )
{
    struct stat info;
    if ( stat(filename, &info) == -1 )          /* cannot stat */
        perror( filename );                    /* say why      */
    else                                         /* else show info */
        show_file_info( filename, &info );
}
```

---

```
void show_file_info( char *filename, struct stat *info_p )
/*
 * display the info about 'filename'. The info is stored in struct at
 *info_p
 */
{
    char    *uid_to_name(), *ctime(), *gid_to_name(), *filemode();
    void    mode_to_letters();
    char    modestr[11];

    mode_to_letters( info_p->st_mode, modestr );

    printf( "%s"      , modestr );
    printf( "%4d "    , (int) info_p->st_nlink);
    printf( "%-8s "   , uid_to_name(info_p->st_uid) );
    printf( "%-8s "   , gid_to_name(info_p->st_gid) );
    printf( "%8ld "   , (long)info_p->st_size);
    printf( "%.12s "  , 4+ctime(&info_p->st_mtime));
    printf( "%s\n"    , filename );
}
```

---

```
/*
 * utility functions
 */

/*
 * This function takes a mode value and a char array
 * and puts into the char array the file type and the
 * nine letters that correspond to the bits in mode.
 * NOTE: It does not code setuid, setgid, and sticky
 * codes
 */
void mode_to_letters( int mode, char str[] )
{
    strcpy( str, "-----" );          /* default=no perms */

    if ( S_ISDIR(mode) ) str[0] = 'd';  /* directory?      */
    if ( S_ISCHR(mode) ) str[0] = 'c';  /* char devices    */
    if ( S_ISBLK(mode) ) str[0] = 'b';  /* block device    */

    if ( mode & S_IRUSR ) str[1] = 'r';  /* 3 bits for user */
    if ( mode & S_IWUSR ) str[2] = 'w';
    if ( mode & S_IXUSR ) str[3] = 'x';

    if ( mode & S_IRGRP ) str[4] = 'r';  /* 3 bits for group */
    if ( mode & S_IWGRP ) str[5] = 'w';
    if ( mode & S_IXGRP ) str[6] = 'x';

    if ( mode & S_IROTH ) str[7] = 'r';  /* 3 bits for other */
    if ( mode & S_IWOTH ) str[8] = 'w';
    if ( mode & S_IXOTH ) str[9] = 'x';
}
```

---

```
#include <pwd.h>

char *uid_to_name( uid_t uid )
/*
 *      returns pointer to username associated with uid, uses getpw()
 */
{
    struct passwd *getpwuid(), *pw_ptr;
    static char numstr[10];

    if ( ( pw_ptr = getpwuid( uid ) ) == NULL ){
        sprintf(numstr,"%d", uid);
        return numstr;
    }
    else
        return pw_ptr->pw_name ;
}

#include <grp.h>
char *gid_to_name( gid_t gid )
/*
 *      returns pointer to group number gid. used getgrgid(3)
 */
{
    struct group *getgrgid(), *grp_ptr;
    static char numstr[10];

    if ( ( grp_ptr = getgrgid(gid) ) == NULL ){
        sprintf(numstr,"%d", gid);
        return numstr;
    }
    else
        return grp_ptr->gr_name;
}
```

# Result

```
$ ./ls2
drwxrwxr-x 4 bruce bruce 1024 Aug 2 18:18 .
drwxrwxr-x 5 bruce bruce 1024 Aug 2 18:14 ..
-rw-rw-r-- 1 bruce users 30720 Aug 1 12:05 s.tar
-rwxrwxr-x 1 bruce users 37351 Aug 1 12:13 tail1
-rw-rw-r-- 2 bruce users 345 Jul 29 11:05 Makefile
-rw-r--r-- 1 bruce users 723 Aug 1 14:26 ls1.c
-rw-r--r-- 1 bruce users 3045 Feb 15 03:51 ls2.c
```

```
$ ls -l
total 189
-rw-rw-r-- 2 bruce users 345 Jul 29 11:05 Makefile
-rw-rw-r-- 1 bruce users 27521 Aug 1 12:14 chap03
drwxrwxr-x 2 bruce users 1024 Aug 1 12:15 docs
-rwxrwxr-x 1 bruce bruce 37048 Aug 1 14:26 ls1
-rw-r--r-- 1 bruce users 723 Aug 1 14:26 ls1.c
-rwxrwxr-x 2 bruce bruce 42295 Aug 2 18:18 ls2
-rw-r--r-- 1 bruce users 3045 Feb 15 03:51 ls2.c
```

} sorting




## Summary

```
readdir() struct dirent {
    ino_t      d_ino;      /* inode number */
    off_t      d_off;      /* not an offset; see NOTES */
    unsigned short d_reclen; /* length of this record */
    unsigned char d_type;   /* type of file; not supported
                             by all filesystem types */
    char        d_name[256]; /* filename */
};
```

```
stat() struct stat {
    dev_t  st_dev;      /* ID of device containing file */
    ino_t   st_ino;      /* inode number */
    mode_t  st_mode;     /* protection */
    nlink_t st_nlink;    /* number of hard links */
    uid_t   st_uid;     /* user ID of owner */
    gid_t   st_gid;     /* group ID of owner */
    dev_t   st_rdev;     /* device ID (if special file) */
    off_t   st_size;     /* total size, in bytes */
    blksize_t st_blksize; /* blocksize for filesystem I/O */
    blkcnt_t st_blocks;  /* number of 512B blocks allocated */
    time_t   st_atime;   /* time of last access */
    time_t   st_mtime;   /* time of last modification */
    time_t   st_ctime;   /* time of last status change */
};
```

```
getpwuid() struct passwd {
    char *pw_name;      /* username */
    char *pw_passwd;     /* user password */
    uid_t pw_uid;       /* user ID */
    gid_t pw_gid;       /* group ID */
    char *pw_gecos;     /* user information */
    char *pw_dir;       /* home directory */
    char *pw_shell;     /* shell program */
};
```

```
getgrgid() struct group {
    char *gr_name;      /* group name */
    char *gr_passwd;    /* group password */
    gid_t gr_gid;       /* group ID */
    char **gr_mem;      /* group members */
};
```



---

# Agenda


- What Does **ls** Do?
- Brief Review of the File System Tree
- How **Dos** **ls** Work?
- Can I Write **ls**?
- Writing **ls -l**
- **Three Special Bits**
- Setting and Modifying the Properties of a File



---

## The Three Special Bits

- The `st_mode` member of the `stat` structure:



- Three special bits are used to activate special properties of a file
  - `suid`(set-user-ID) bit
  - `sgid`(set-group-ID) bit
  - sticky bit

---

# 1. The Set-User-ID Bit

- How can a regular user change his or her password?
  - Use the passwd command!
  - But, how does the passwd command work?

```
$ ls -l /etc/passwd
-rw-r--r--  1 root    root          894 Jun 20 19:17 /etc/passwd
```

## Problem:

Changing your password means changing your record in the file `/etc/passwd`, but you do **NOT** have **permission** to write to that file.

Only the user named **root** has write permission.

---

# 1. The Set-User-ID Bit

- Solution: Give permission to the program, not to you.

```
$ ls -l /usr/bin/passwd  
-r-sr-xr-x  1 root    bin          15725 Oct 31  1997 /usr/bin/passwd
```

- The program you use to change your password, /usr/bin/passwd or /bin/passwd, is owned by root and has the set-user-ID (SUID ) bit set.
- That **SUID** bit tells the kernel to run the program as though it were being run by the owner of the program.

---

# 1. The Set-User-ID Bit

- Doesn't that mean I can change passwords of other users?
  - NO;
    - The `passwd` program knows who you are.
    - It uses the `getuid` system call to ask the kernel for the user ID you used when you logged in.
  - `passwd` has permission to rewrite the entire password file, but will **ONLY** change the **record for the user** running the program.
- Program can test whether a file has SUID bit on by using the mask defined in `<sys/stat.h>`

```
#define S_ISUID          0004000          /* set user id on execution */
```

---

## 2. The Set-Group-ID Bit

- The SGID bit sets the effective group ID of a program
  - If a program belongs to group *g* and the set-group-ID bit is set, the program runs as though it were being run by a member of group *g*
- This bit grants the program the access rights of members of that group
- A mask to test for the SGID bit

```
#define S_ISGID          0002000          /* set group id on execution */
```

---

## 3. The Sticky Bit

### ■ Use for files

- In swapping, the sticky bit told the kernel to keep the program on the swap device so that kernel can load it faster.
  - Loading program from swap device is fast because program was never fragmented on the swap device.
  - Now, **no longer necessary** due to virtual memory and paging that allow the kernel to move programs in and out of memory **in small sections**.

### ■ Use for directories

- /tmp are publicly writable, allowing any user to create and delete any files there.
- The sticky bit overrides the publicly writable attribute for a directory.  
**Files in the directory may ONLY be deleted by their owners if the sticky bit is set**

---

## The Special Bits and ls -l

- Each file has a type and 12 attribute bits, but **ls** uses only 9 spots to display these 12 attributes.

```
-rwsr-sr-t  1 root      root          2345 Jun 12 14:02 sample
```

- Letter **s** indicates that the user and group-executable bits have been augmented by the set-user and set-group ID bits.
- Letter **t** at the end indicates that the sticky bits is on.

---

# Agenda

- What Does **ls** Do?
- Brief Review of the File System Tree
- How **Dos ls** Work?
- Can I Write **ls**?
- Writing **ls -l**
- Three Special Bits
- Setting and Modifying the Properties of a File



---

## Type of a File


```
drwxr-xr-x 2 root root 0 Jan 1 1970 home
```

- A file has a type
  - It can be a regular file(-), a directory(d), a device file(b, c), a socket(s), a symbolic link(l), or a named pipe(p).
  
- The type of the file is established when the file is created.
  - The creat() system call creates a regular file.
  - Different system call are used to create directories and devices.
  
- It is not possible to change the type of a file.

---

## Permission Bits and Special Bits

- Every file has 9 permission bits and 3 special bits.



- These bits are established when file is created and can be modified by making the chmod system call

```
fd = creat( "newfile", 0744 );
```

- If you want to prevent programs from creating files that can be modified by group or others
  - `umask(022);`

---

# Permission Bits and Special Bits

- Changing the mode of a file: **chmod()** system call

```
chmod( "/tmp/myfile", 04764 );
```

or

```
chmod( "/tmp/myfile", S_ISUID | S_IRWXU | S_IRGRP|S_IWGRP | S_IROTH );
```

- A shell command to change permission and special bits

```
$ chmod 04764 test
```

or

```
$ chmod u=rws test
```

```
$ chmod g=rw test
```

```
$ chmod o=r test
```

---

## **chmod**

---

<b>PURPOSE</b>	Change permission and special bits for a file	
<b>INCLUDE</b>	#include <sys/types.h> #include <sys/stat.h>	
<b>USAGE</b>	int result = chmod(char *path, mode_t mode);	
<b>ARGS</b>	path	path to file
	mode	new value for mode
<b>RETURNS</b>	-1	if error
	0	if success

---

---

## Number of Links to a File

- The number of links is simply the number of times the file is referenced in directories.
  - If a file appears in three places in various directories, the link count is 3. (in the next chapter)

```
$ ls -l
total 108
-rw-rw-r-- 2 bruce  users      345 Jul 29 11:05 Makefile
-rw-rw-r-- 1 bruce  users    27521 Aug  1 12:14 chap03
drwxrwxr-x 2 bruce  users     1024 Aug  1 12:15 docs
Type&permission  links  owner    group    size    modified-date/time  name
```

---

## Owner and Group of a File

- Establishing the owner of a file:
  - The owner of file is the user who creates it
  - When kernel creates a file, it sets the owner of the file to be the effective user ID of the process that calls `creat()`
  - If the program has the set-user-ID bit set, though, the effective user ID is the user ID of the person who owns the program.

---

# Owner and Group of a File

- Establishing the group of a file:
  - The group of a file is set to the effective group ID of the process that creates the file.
  - Under non-ordinary circumstances, the group ID of a file is set to the group ID of the parent directory.

---

# Owner and Group of a File

## ■ Changing the owner and group of a File

### ○ **chown()** system call:

- Normally, users do not change the owner of a file
- Typically used to set up and manage user accounts

```
chown( "file1", 200, 40 );
```

## ■ Shell Commands to Change User and Group ID for Files: **chown, chgrp**

```
[seokin@compasslab2 ch03]$ ls -al ls2
-rwxrwxr-x. 1 seokin seokin 13224 Sep 16 09:48 ls2
[seokin@compasslab2 ch03]$ sudo chown jhong ls2
[seokin@compasslab2 ch03]$ ls -al ls2
-rwxrwxr-x. 1 jhong seokin 13224 Sep 16 09:48 ls2
```

```
[seokin@compasslab2 ch03]$ sudo chgrp jhong ls2
[seokin@compasslab2 ch03]$ ls -al ls2
-rwxrwxr-x. 1 jhong jhong 13224 Sep 16 09:48 ls2
```



---

## **chown**

---

**PURPOSE**      Change owner and or group ID of a file

---

**INCLUDE**      #include <unistd.h>

---

**USAGE**          int chown(char \*path, uid\_t owner, gid\_t group)

---

<b>ARGS</b>	path	path to file
	owner	user ID for file
	group	group ID for file

---

<b>RETURNS</b>	-1	if error
	0	if success

---

---

# Modification and Access Time

- Each file has three timestamps of
  - last modified
  - last read
  - file properties (such as owner ID or permission bits) were last changed
  - Kernel automatically updates these times as programs read and write the file
- Changing modification and access times of a file:
  - **utime()** system call
- Shell Commands: **touch**

---

## **utime**

---

<b>PURPOSE</b>	Change access and modification time for files	
<b>INCLUDE</b>	#include <sys/time.h> #include <utime.h>	
<b>USAGE</b>	#include <sys/types.h> int utime( char *path, struct utimbuf *newtimes )	
<b>ARGS</b>	path	path to file
	newtimes	pointer to a struct utimbuf see utime.h for details
<b>RETURNS</b>	-1	if error
	0	if success

---

---

## Name of a File

- Establishing the Name of a File

- `creat()` system call sets the name and the initial mode of a file.

- Changing the Name of a File:

- `rename()` system call

- Shell Command : `mv`

- Allows you to change the name of a file
- Also allows you to move a file from one directory to another

---

## **rename**

---

<b>PURPOSE</b>	Change name and/or move a file
<b>INCLUDE</b>	#include <stdio.h>
<b>USAGE</b>	int result = rename( char *old, char *new )
<b>ARGS</b>	old    old name of file or directory new    new pathname for file or directory
<b>RETURNS</b>	-1    if error 0     if success

---

---

# Objectives

## ■ Ideas and Skills

- A directory is a list of files
- How to read a directory
- Types of files and how to determine the type of a file
- Properties of files and how to determine properties of a file
- Bit sets and bit masks
- User and group ID numbers

## ■ System Calls and Functions

- opendir, readdir, closedir, seekdir
- stat
- chmod, chown, utime
- rename

## ■ Commands

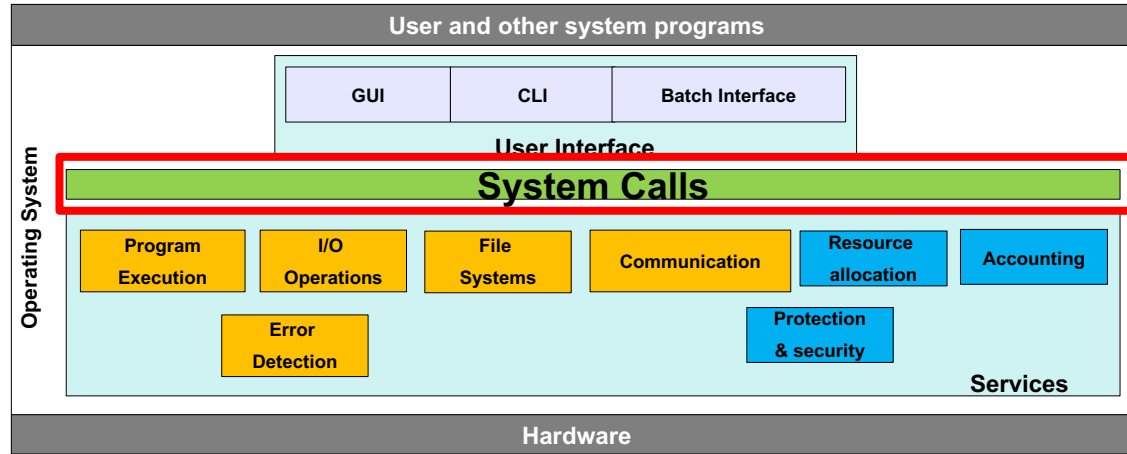
- ls

---

# Next?

- **Ch04. File System**

# System Calls



- **System call** : Programmatic way in which a computer program requests a service of the operating system.
- **A function** provided to applications by the OS kernel
  - Generally to use a hardware abstraction (file, socket)
  - Or to use OS-provided software abstraction (IPC, scheduling)
- System calls are the **only entry points** into the kernel system
  - All programs needing resources must use system calls



## System Calls

- Why not put these directly in the application?
  - **Protection of OS/hardware from buggy/malicious programs**
  - Applications are not allowed to directly interact with hardware, or access kernel data structures
  - OS must validate system call parameters

# Interrupt


## ■ Interrupt view of CPU

```
while (fetch next instruction){  
    run instruction;  
    if(there is an interrupt){  
        save CPU context and error code if any  
        find OS-provided interrupt handler  
        jump to handler  
        restore CPU context when handler returns  
    }  
}
```

# Software Interrupt

- “**int <num>**” instruction
  - allows software to raise an interrupt (**software Interrupt**)
  - <num> is 0x80 in Linux

## x86 interrupt table



## How to invoke System Calls?

1. Kernel assigns a **system call number** to each system call type and initialize the **system call table**
- 2. User process sets up system call number and arguments**
- 3. User process runs int X (X is 0x80 in linux0)**
4. Hardware **switches to kernel mode** and invokes kernel's **interrupt handler** for X (interrupt dispatch)
5. Kernel looks up **system call table** using system call number
6. Kernel **invokes the corresponding function**
7. Kernel returns by running iret (interrupt return)

# How to invoke System Calls?

## Example

```
.data
    s:
        .ascii "hello world\n"
        len = . - s
.text
    .global _start
    _start:

    movl $4, %eax    /* write system call number */
    movl $1, %ebx    /* stdout */
    movl $s, %ecx    /* the data to print */
    movl $len, %edx /* length of the buffer */
    int $0x80

    movl $1, %eax    /* exit system call number */
    movl $0, %ebx    /* exit status */
    int $0x80
```

```
as -o main.o main.S
ld -o main.out main.o
./main.out
```

## Why use interrupts to invoke system call?

- Also protection
- Forces applications to call well-defined “public functions”
- **But, using interrupt is complex and not portable!**


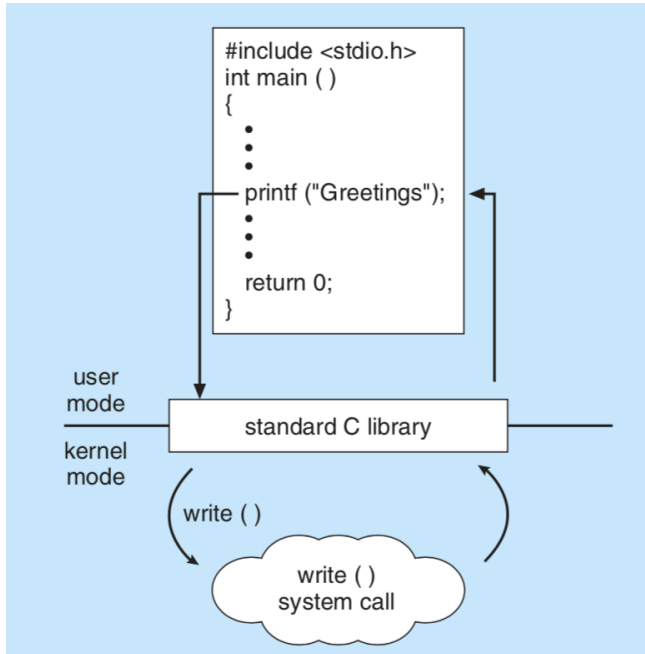
## API for System Calls

- Programmers use system calls indirectly through **APIs** (Application Programming Interface: set of functions)
  - **Windows API** for Windows systems
  - **POSIX API** for POSIX-based systems (UNIX, Linux, and Mac OS X)
  - **Java API** for programs that run on the JVM
- **A programmer accesses an API via a library**
  - Ex) libc

<u>glibc</u>	<u>system call</u>
write	write
read	read
printf	write
fread	read
malloc	brk
pthread_lock	futex

# System Call Interface

- Set of library functions that links to the system calls



- Caller does not need to know how system call is implemented
- Caller needs to know only the interface and what it returns

```
ssize_t write(int fd, void *buf, size_t count);
```



## System Call Types

### ■ Six categories

#### ○ Process control

- fork, exec, exit ...

#### ○ File manipulation

- create, open, close, read, write, lseek

#### ○ Device manipulation

- open, close, read, write, ioctl

#### ○ Information maintenance

- time, date, dump, pid

#### ○ Communications

- open, close, connect, accept, read, write, send, recv, pipe, mmap, sendfile  
...

#### ○ Protection

- chmod, umask, chown ...