

Ch10-1. I/O Redirection

Prof. Seokin Hong

Kyungpook National University

Fall 2019

THE SHELL

- A shell is a program that manages processes and runs programs
- Thee main functions of shells
 - (a) Shells run programs
 - (b) Shells manage input and output
 - (c) Shells can be programmed

THE SHELL

- How do the following commands work?

```
$ ls > my.files
```

```
$ who | sort > userlist
```

- We focus on a particular form of inter-process communication: **I/O redirection and pipes**

Objectives

■ Ideas and Skills

- I/O Redirection : What and why?
- Definitions of standard input, output, and error
- Redirecting standard I/O to files
- Using fork to redirect I/O for other programs
- **Pipes**
- Using fork with pipes

■ System Calls and Functions

- dup, dup2
- pipe

A SHELL APPLICATION : WATCH FOR USERS

- Consider the following problem:
You want a program that notifies you when people log in or log out of the system.
- You could write a C program that uses the utmp file
- A simpler solution is to write a shell script : who

A SHELL APPLICATION : WATCH FOR USERS

Logic

Get list of users (call it prev)

While true

 sleep

 get list of users (call it curr)

 compare lists

 in prev, not in curr -> logout

 in curr, not in prev -> login

 make prev = curr

repeat

Shell code

Who | sort > prev

While true ; do

 sleep 60

 who | sort > curr

 echo "logged out:"

 comm -23 prev curr

 echo "logged in:"

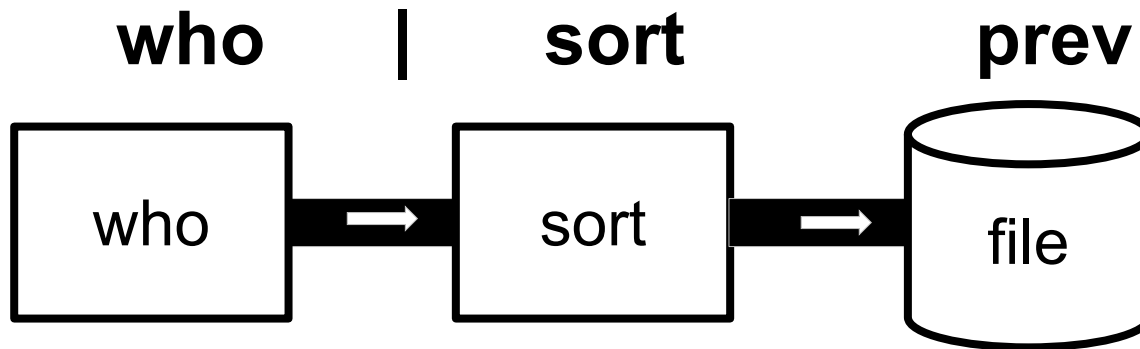
 comm -13 prev curr

 mv curr prev

done

A SHELL APPLICATION : WATCH FOR USERS

- `who | sort > prev`



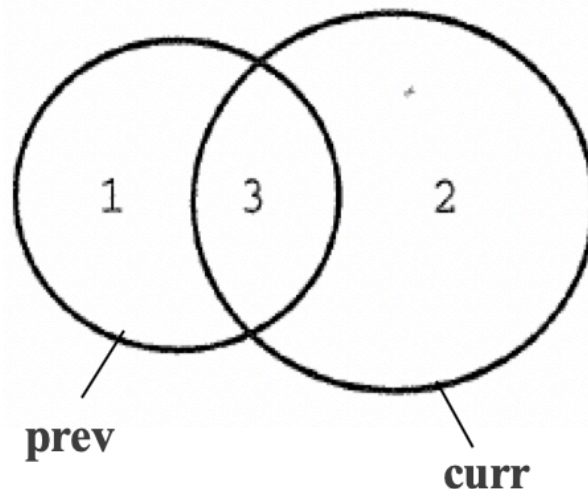
shell code

```
-----  
who | sort > prev  
while true ; do  
    sleep 60  
    who | sort > curr  
    echo "logged out:"  
    comm -23 prev curr  
    echo "logged in:"  
    comm -13 prev curr  
    mv curr prev  
done
```

A SHELL APPLICATION : WATCH FOR USERS

■ The **comm** command

- compares two sorted lists
- **comm** : print out three columns : 1,2,3
- **comm -23 prev curr**
 - drop columns 2 and 3 → show lines only in prev
- **comm -13 prev curr**
 - drop columns 1 and 3 → show lines only in curr



shell code

```
-----  
who | sort > prev  
while true ; do  
    sleep 60  
    who | sort > curr  
    echo "logged out:"  
    comm -23 prev curr  
    echo "logged in:"  
    comm -13 prev curr  
    mv curr prev  
done
```


A SHELL APPLICATION : WATCH FOR USERS

■ Lessons:

- The | operator passes the output of one command as input to another
- > directs the output of a command into a file.

`x = func_a(func_b(y));` in C

`prog_b | prog_a > x` in sh

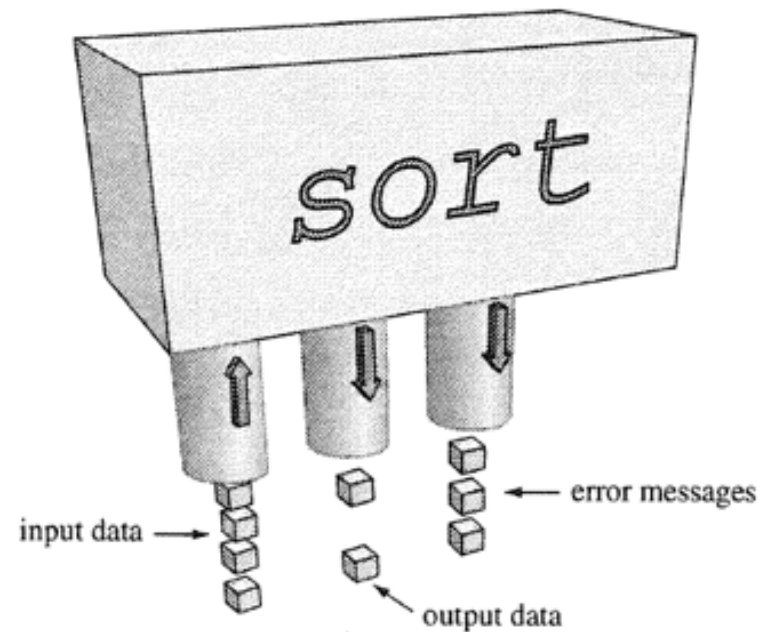
■ Question: How?

Contents

- 10.1 Shell
- 10.2 A Shell Application : Watch for Users
- 10.3 Facts about Standard I/O and Redirection
- 10.4 How to Attach stdin to a File
- 10.5 Redirecting I/O for Another Program: `who>userlist`
- 10.6 Programming Pipes

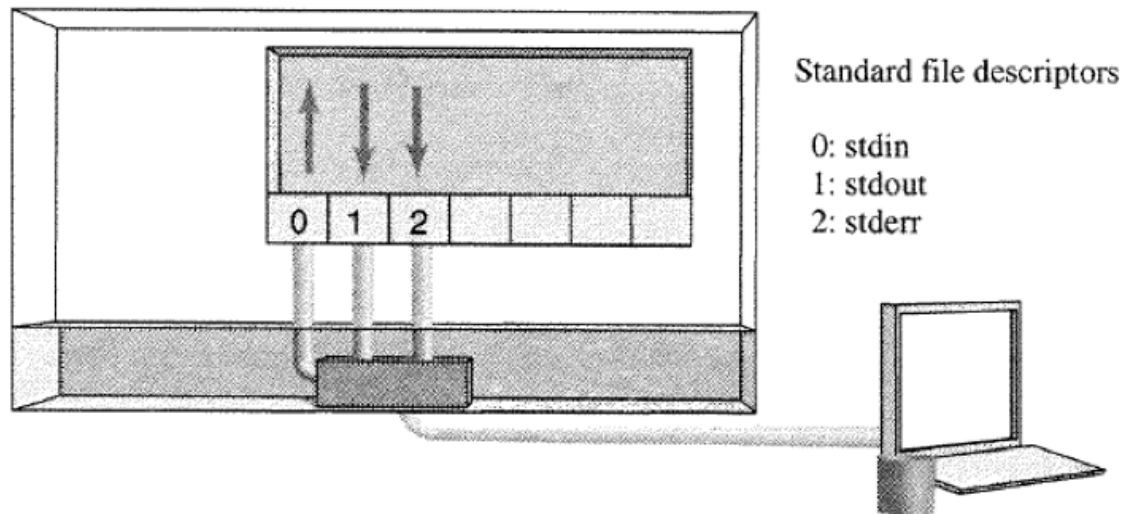
FACT ABOUT STANDARD I/O AND REDIRECTION

- All unix I/O redirection is based on the principle of **standard streams of data**.
- **Standard input**
 - the stream of data to process
- **Standard output**
 - the stream of result data
- **Standard error**
 - a stream of error messages



Fact One: Three Standard File Descriptors

- All Unix/Linux commands use file descriptor 0, 1, and 2.
 - 0: standard in (stdin)
 - 1: standard out (stdout)
 - 2: standard error (stderr)
- Every commands (programs) gets three open file descriptors at startup:



Output Goes Only to stdout

- Most programs do NOT accept names for output files;
 - They always write **results** to **file descriptor 1** and **errors** to **file descriptor 2**.
 - If you want to send the output of a process to a file or to the input of another process, you need to change where the file descriptor goes.

The Shell, Not the Program, Redirects I/O

- You **tell the shell to attach file descriptor 1 to a file** by using the output redirection notation:
 \$ cmd > filename
- The shell **connects that file descriptor to the named file.**
- The **program continues to write to file descriptor 1, unaware of the new data destination.**

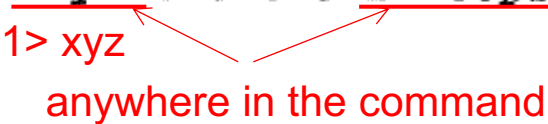
Ex1. listargs.c

```
/* listargs.c
 *          print the number of command line args, list the args,
 *          then print a message to stderr
 */
#include    <stdio.h>

main( int ac, char *av[] )
{
    int     i;

    printf("Number of args: %d, Args are:\n", ac);
    for(i=0;i<ac;i++)
        printf("args[%d] %s\n", i, av[i]);           // to stdout
    fprintf(stderr, "This message is sent to stderr.\n"); // to stderr
}
```

```
$ cc listargs.c -o listargs
$ ./listargs testing one two
args[0] ./listargs
args[1] testing
args[2] one
args[3] two
This message is sent to stderr.
$ ./listargs testing one two > xyz
This message is sent to stderr.
$ cat xyz
args[0] ./listargs
args[1] testing
args[2] one
args[3] two
$ ./listargs testing >xyz one two 2> oops
$ cat xyz
args[0] ./listargs
args[1] testing
args[2] one
args[3] two
$ cat oops
This message is sent to stderr.
```



1> xyz
anywhere in the command

Understanding I/O Redirection

■ Goal:

- Understand how I/O redirection works
- Learn how to write programs that use it

who > userlist

attach stdout to a file

sort < data

attach stdin to a file

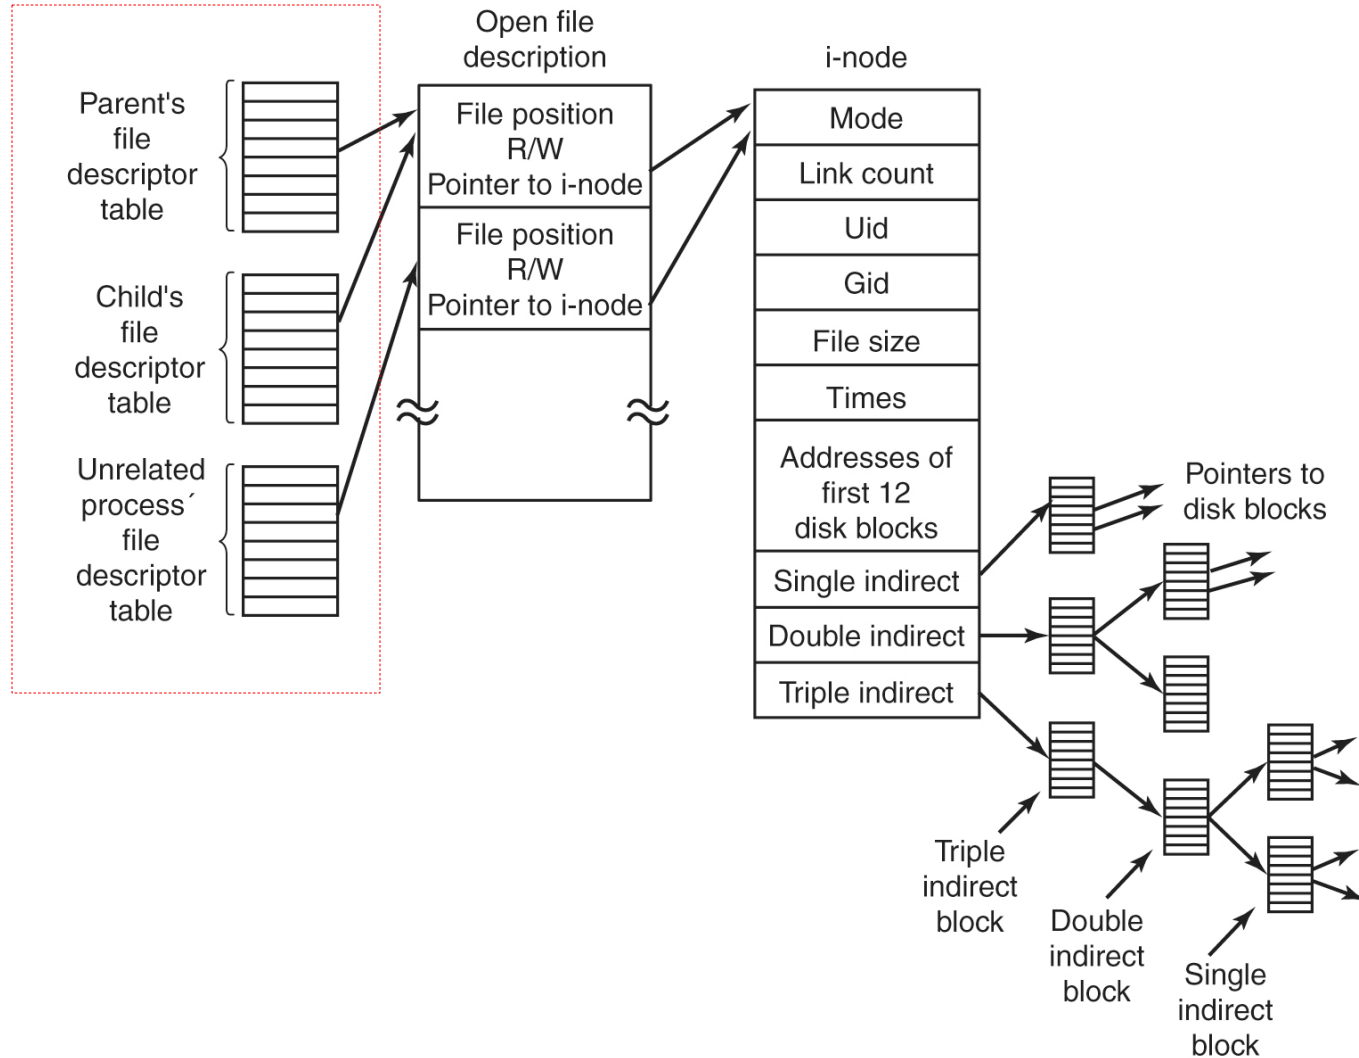
who | sort

attach stdout to stdin

Fact Two: The “Lowest-Available-fd” Principle

- What is a **file descriptor**? It is an index of an array
 - Each process has a collection of files it has open
 - The information of those open files are kept in an array
- **FACT:** When you open a file, you always get the lowest available spot in the array.

File descriptor



Contents

- 10.1 Shell
- 10.2 A Shell Application : Watch for Users
- 10.3 Facts about Standard I/O and Redirection
- 10.4 How to Attach stdin to a File
- 10.5 Redirecting I/O for Another Program: `who>userlist`
- 10.6 Programming Pipes

How to Attach stdin to a File

- How does a program redirect standard input so that data come from a file?

ex) \$ sort < data

```
[Seokinui-MacBookPro:KNU seokin$ cat test
```

```
5
```

```
4
```

```
3
```

```
6
```

```
9
```

```
1
```

```
[Seokinui-MacBookPro:KNU seokin$ sort < test
```

```
1
```

```
3
```

```
4
```

```
5
```

```
6
```

```
9
```

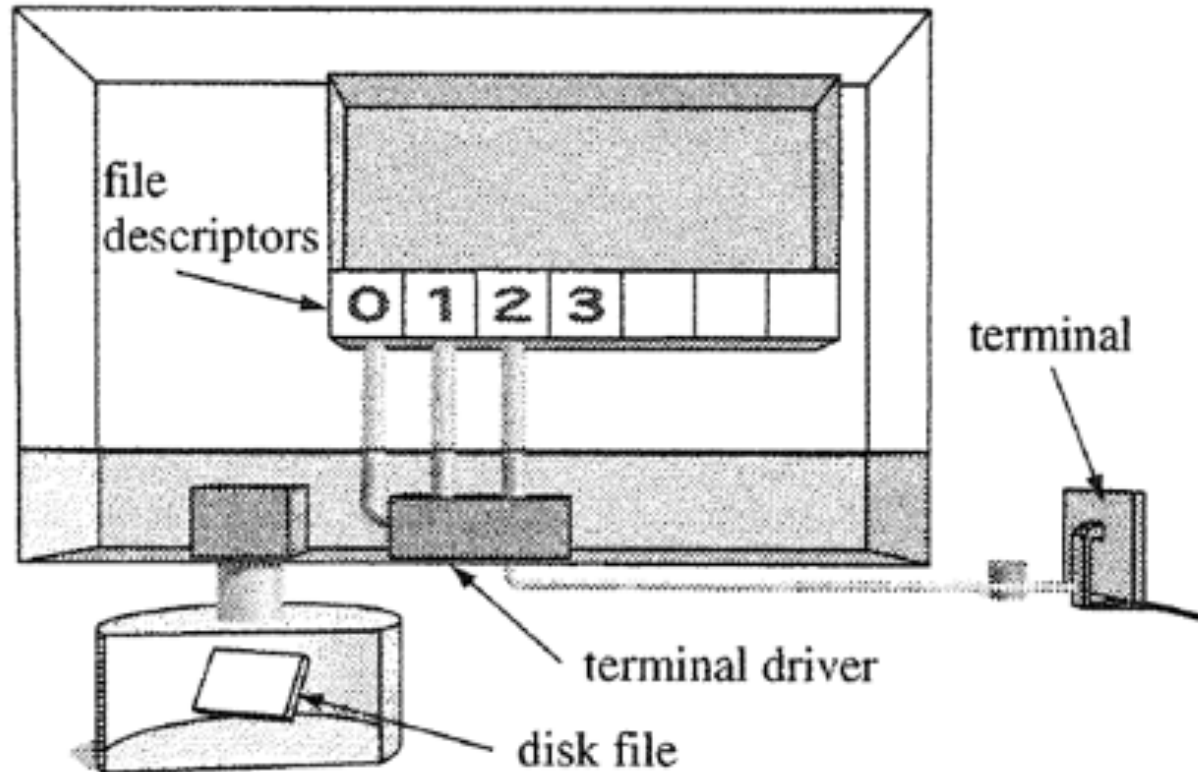
- If we attach file descriptor 0 to a file, that file becomes the source for standard input.

- How? ...

Method 1: Close Then Open

■ Starting

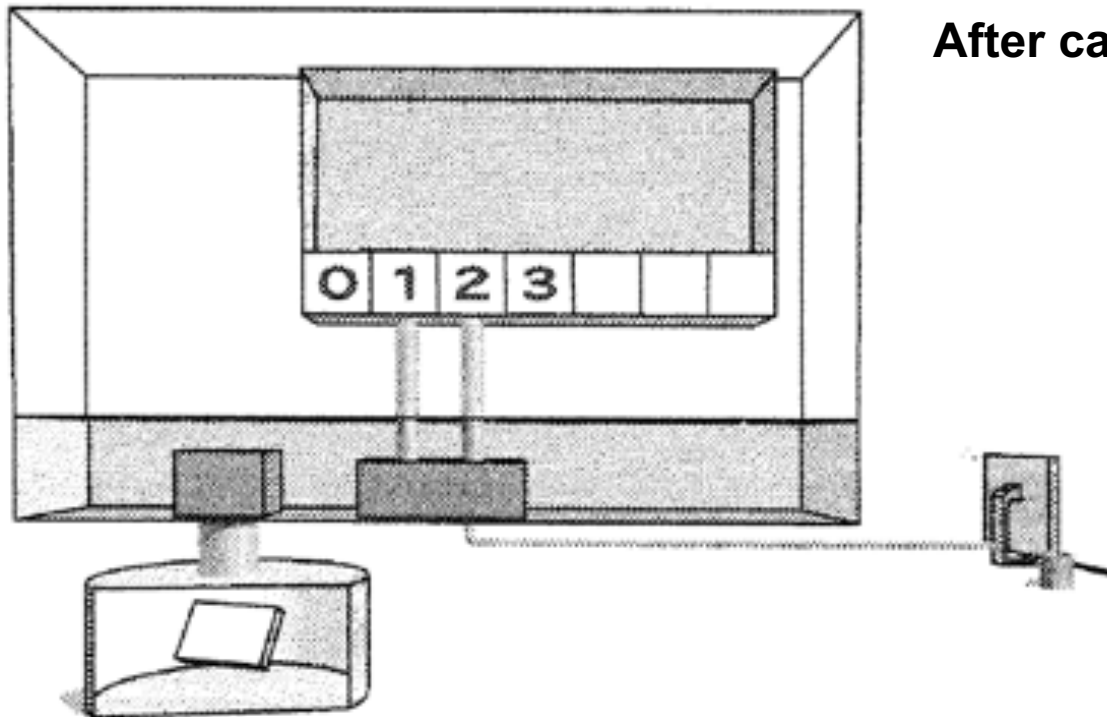
- File descriptor 0, 1, 2 attached to the terminal driver



Method 1: Close Then Open

- **Then, close(0)**

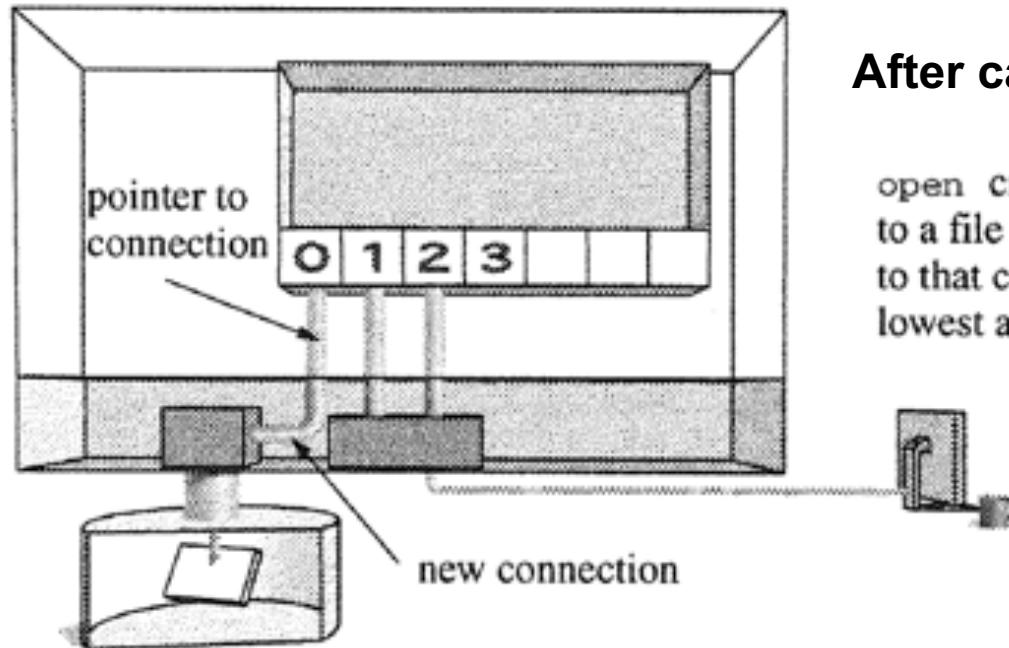
- The first element in the array of file descriptors is now unused



Method 1: Close Then Open

■ Finally, `open(filename, O_RDONLY)`

- Opens the file you want to attach to stdin.



After calling `open()`

`open` creates a connection to a file and puts a pointer to that connection in the lowest available entry.

Ex2. stdinredir1.c

```
#include <stdio.h>
#include <fcntl.h>
#include <stdlib.h>
#include <unistd.h>

void main(void)
{
    int fd;
    char line[100];

    fgets(line, 100, stdin); printf("%s", line);
    fgets(line, 100, stdin); printf("%s", line);
    fgets(line, 100, stdin); printf("%s", line);

    close(0);
    fd = open("/etc/passwd", O_RDONLY);
    if( fd != 0 )
    {
        fprintf(stderr, "Could not open data as fd()\n");
        exit(1);
    }
    fgets(line, 100, stdin); printf("%s", line);
    fgets(line, 100, stdin); printf("%s", line);
    fgets(line, 100, stdin); printf("%s", line);
}
```

Method 2: open..**dup**..close

Method 3: open..**dup2**..close

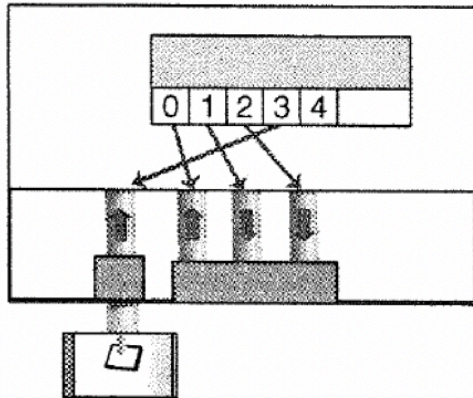
dup, dup2	
PURPOSE	Copy a file descriptor
INCLUDE	#include <unistd.h>
USAGE	newfd = dup(oldfd); newfd = dup2(oldfd, newfd);
ARGS	oldfd file descriptor to copy newfd copy of oldfd
RETURNS	-1 if error newfd new file descriptor

```
#ifdef CLOSE_DUP
    close(0);
    newfd = dup(fd);
#else
    newfd = dup2(fd, 0);
#endif
```

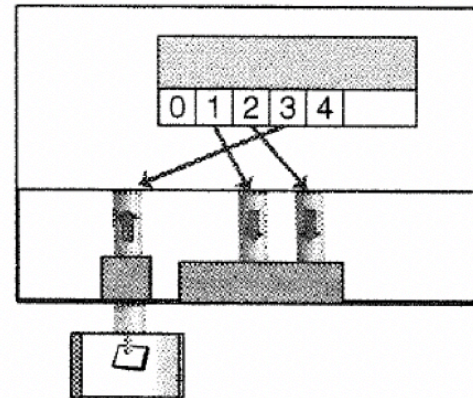
Using dup to redirect

- `dup()` makes a duplication of `fd`
- The duplicate uses the lowest unused file descriptor

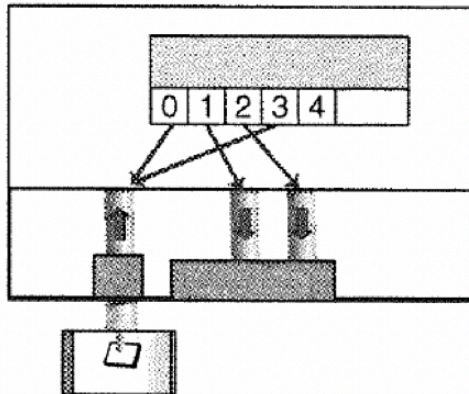
1. `fd = open("f", O_RDONLY);`



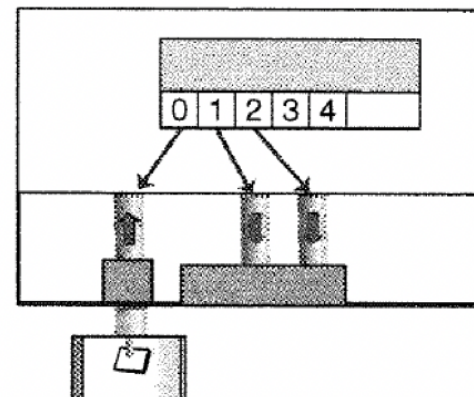
2. `close(0);`



3. `dup(fd);`



4. `close(fd);`



Ex3. stdinredir2.c

```
#include <stdio.h>
#include <fcntl.h>
#include <stdlib.h>
#include <unistd.h>

#define CLOSE_DUP          /* open, close, dup, close */
/*#define USE_DUP2        /* open, dup2, close */

void main(void)
{
    int fd;
    int newfd;
    char line[100];

    fgets(line, 100, stdin); printf("%s", line);
    fgets(line, 100, stdin); printf("%s", line);
    fgets(line, 100, stdin); printf("%s", line);

    fd = open("/etc/passwd", O_RDONLY);

#ifdef CLOSE_DUP
    close(0);
    newfd = dup(fd);
#else
    newfd = dup2(fd, 0);
#endif
    if(newfd != 0){
        fprintf(stderr, "Could not duplicate fd to 0\n");
        exit(1);
    }
    close(fd);

    fgets(line, 100, stdin); printf("%s", line);
    fgets(line, 100, stdin); printf("%s", line);
    fgets(line, 100, stdin); printf("%s", line);
}
```

But the Shell Redirects stdin for Other Programs

- In practice, of course, if a program wants to read a file, it can just open the file directly rather than changing standard input
- The real value of these samples is to show **how one program can change standard input for another program;**
\$ sort < data

Contents

- 10.1 Shell Programming
- 10.2 A Shell Application : Watch for Users
- 10.3 Facts about Standard I/O and Redirection
- 10.4 How to Attach stdin to a File
- 10.5 Redirecting I/O for Another Program: `who > userlist`
- 10.6 Programming Pipes

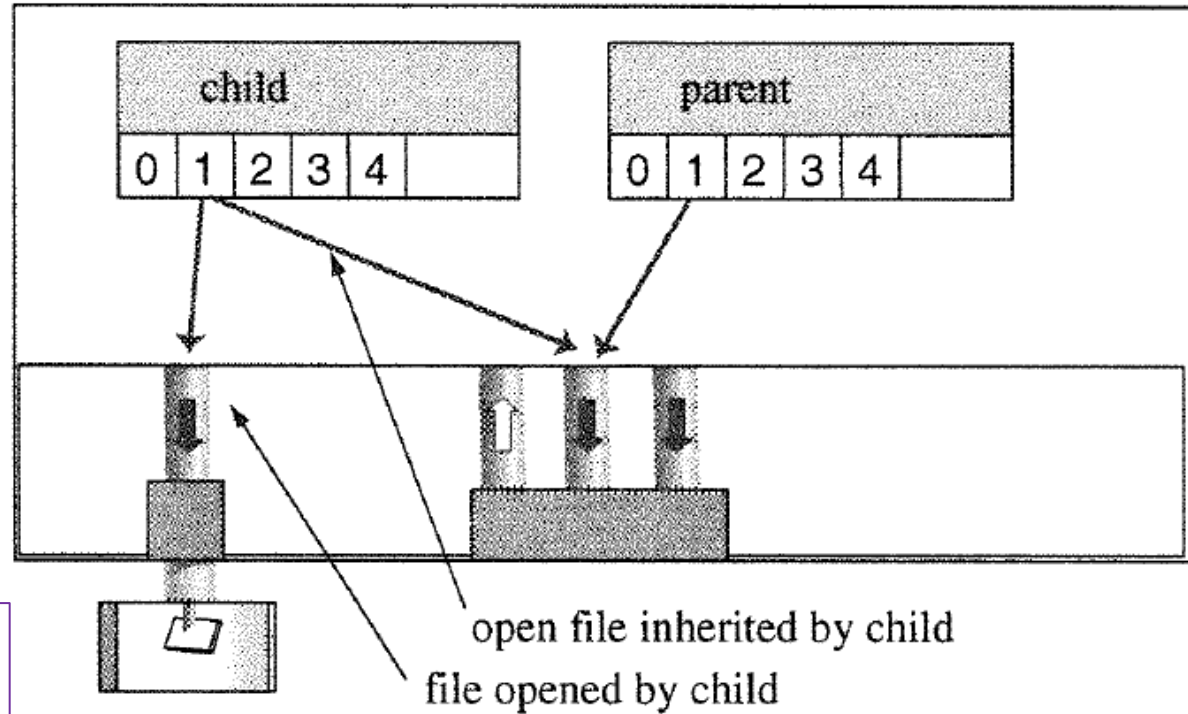
Redirecting I/O for other programs

- The shell redirects output for a child.
 - How “who > userlist” works?

The child inherits from the parent the pointers to open files. The child redirects standard output:

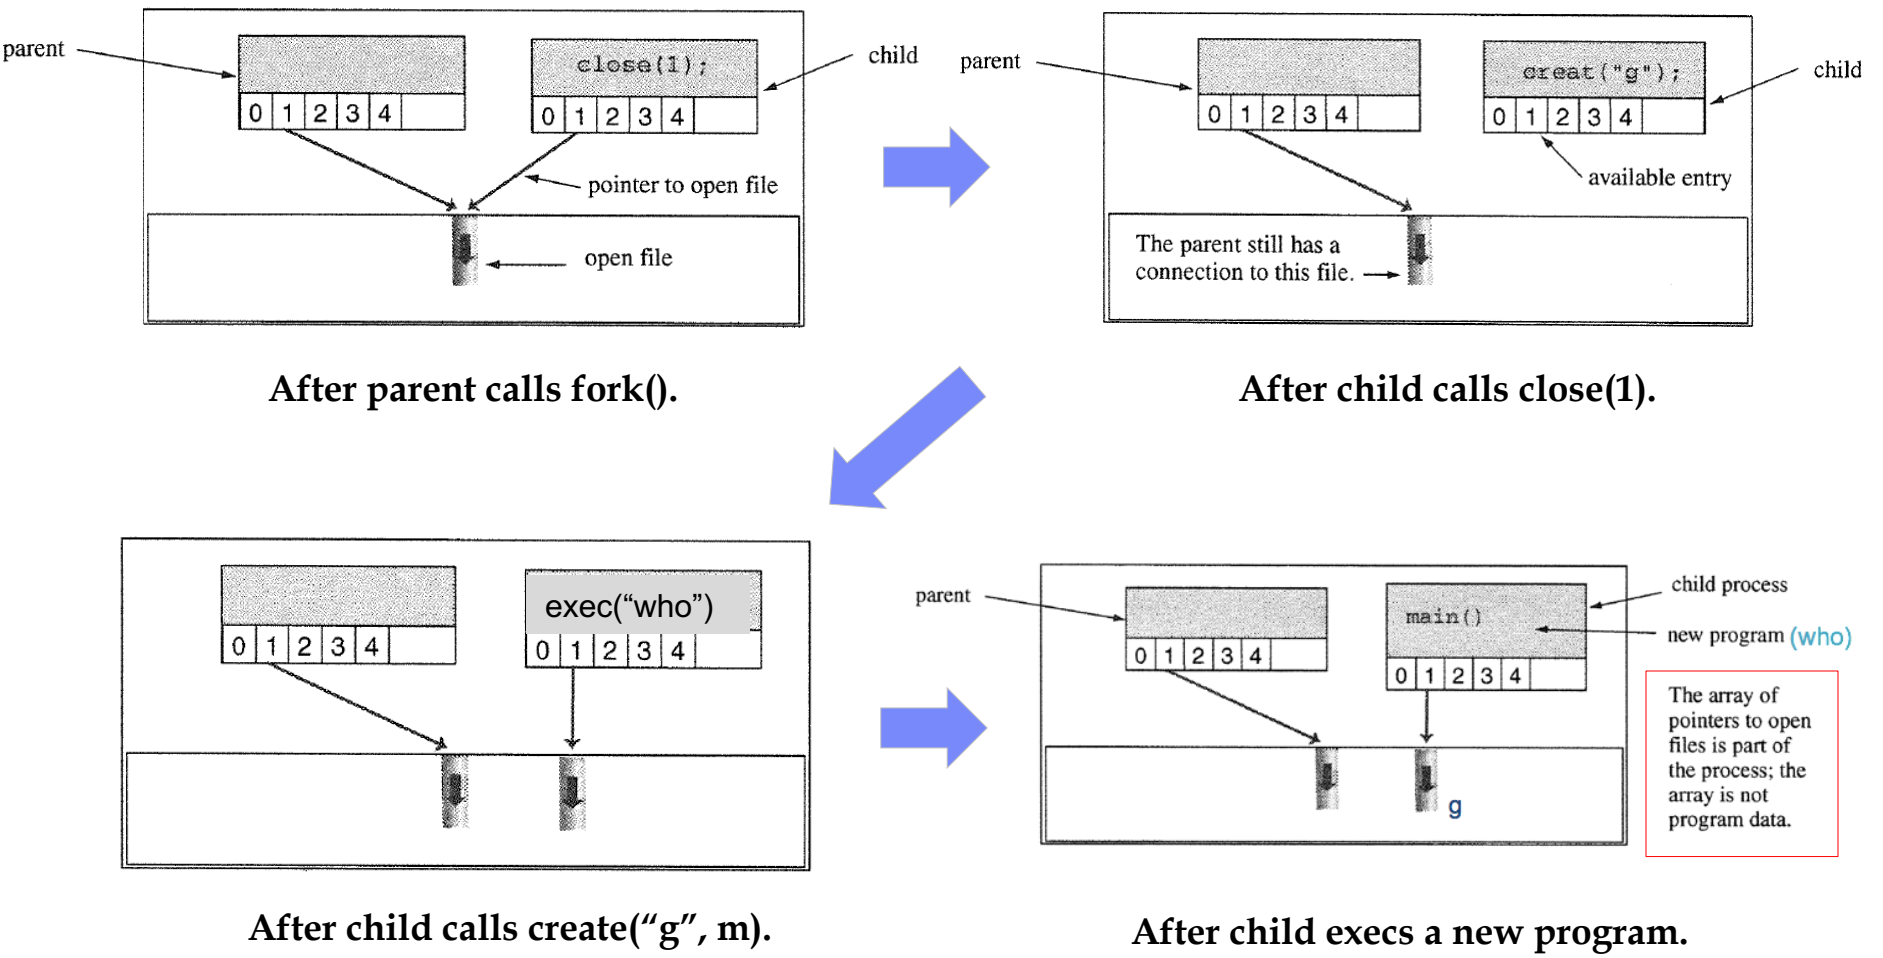
```
close(1);  
creat("f");  
exec();
```

```
creat("userlist");  
exec("who");
```

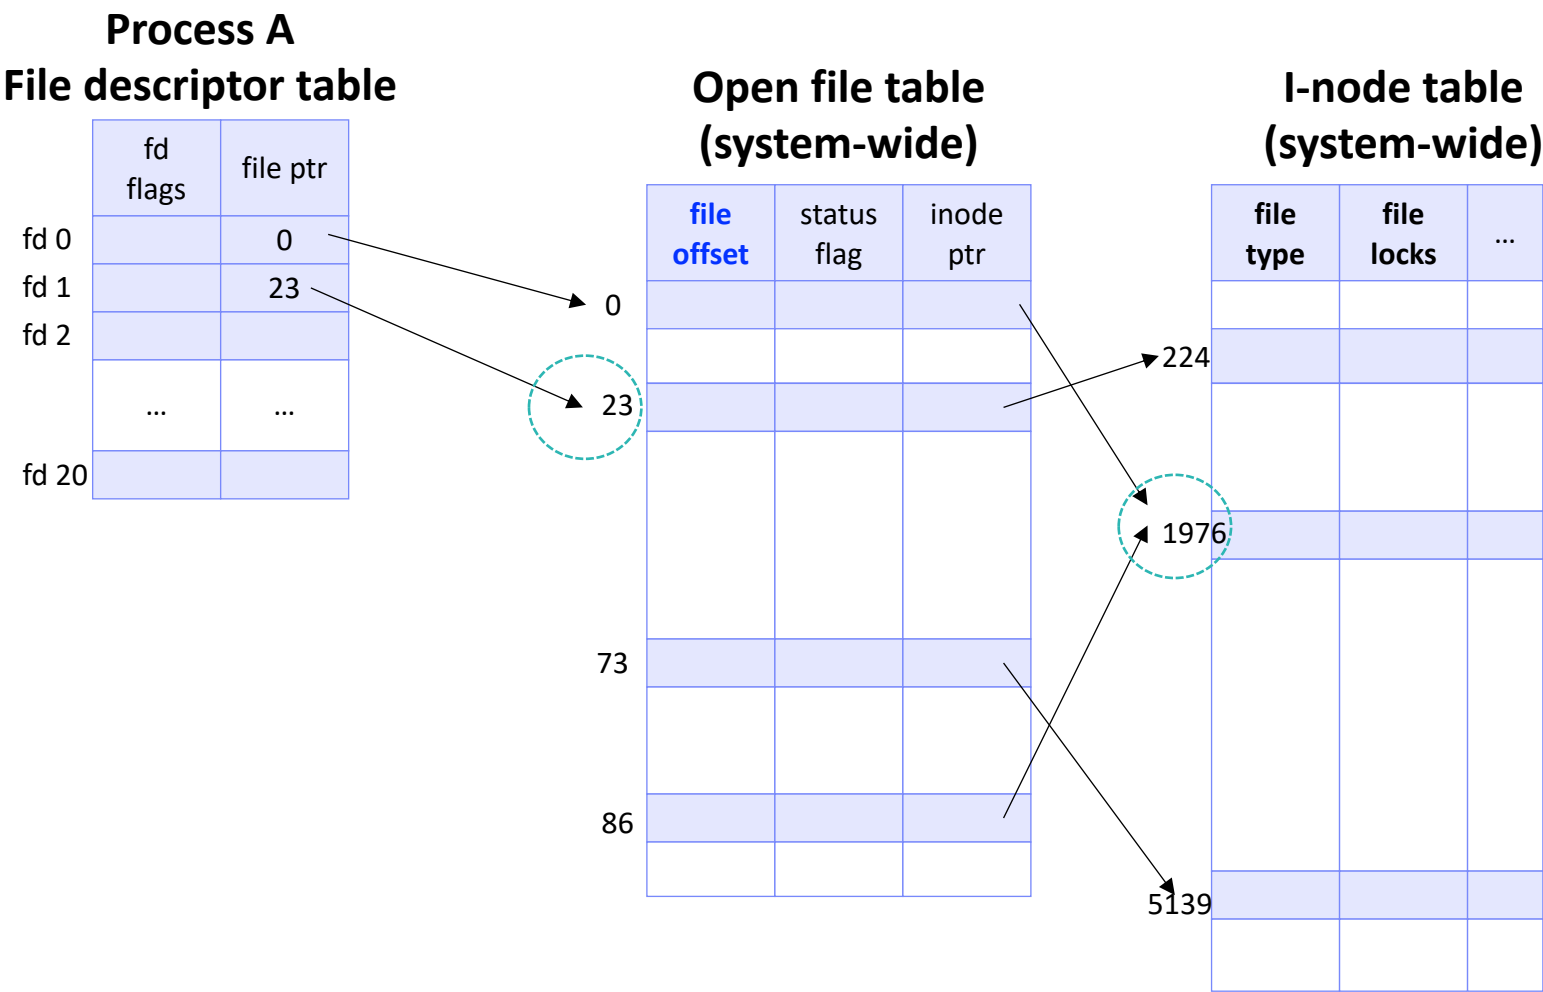


Redirecting I/O for other programs

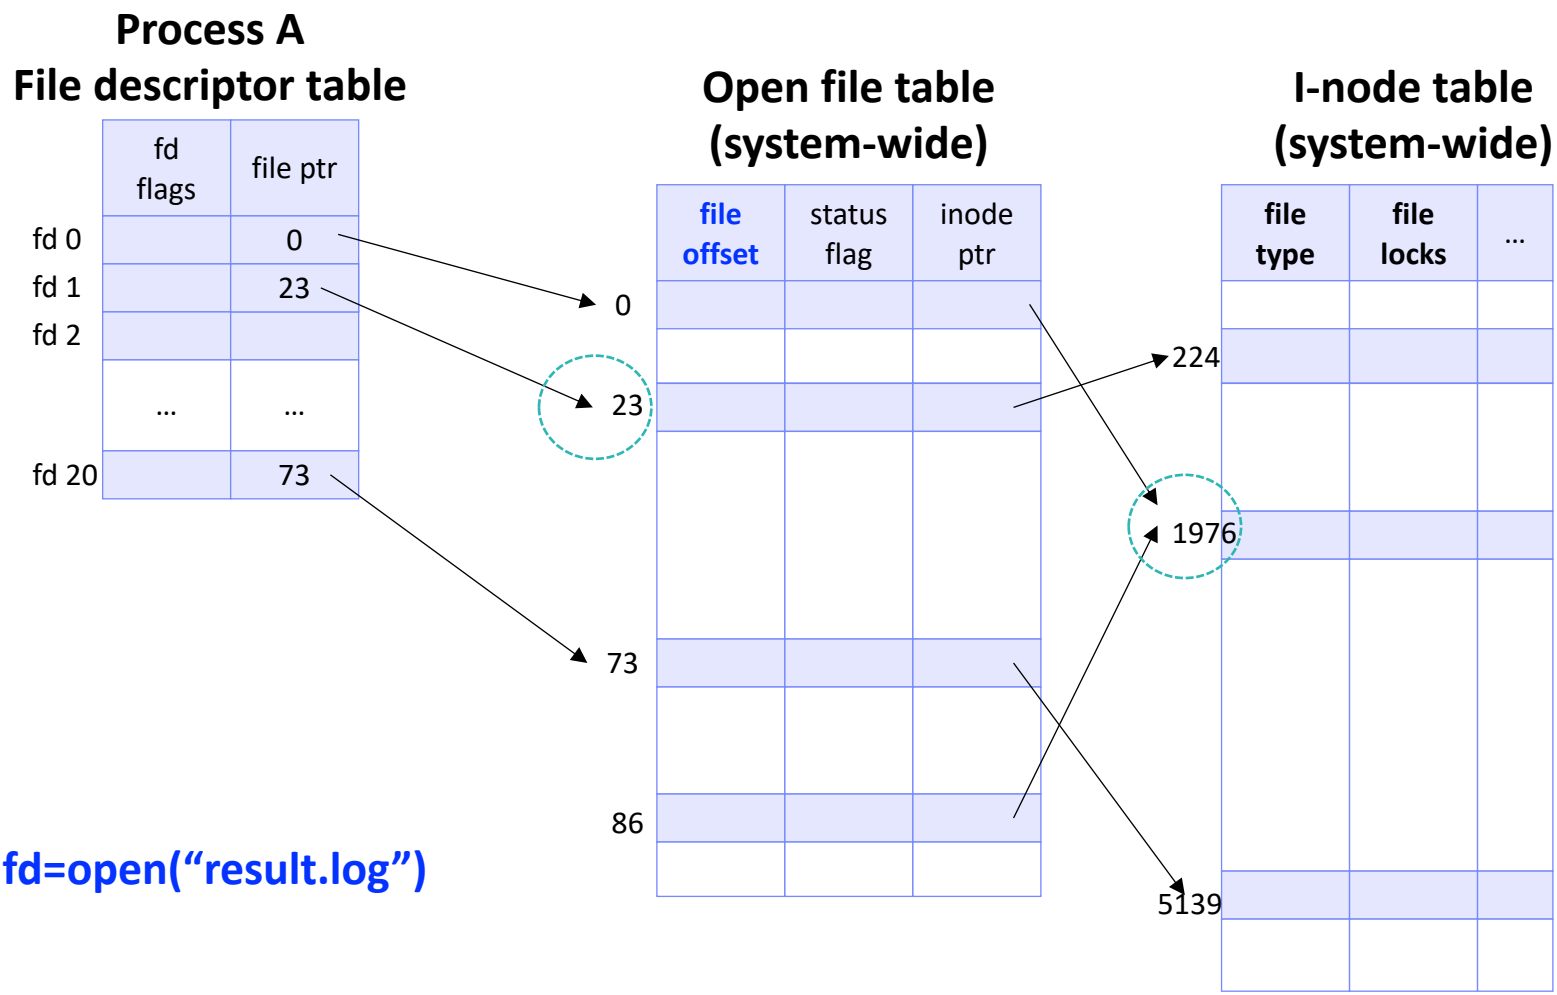
■ who > g



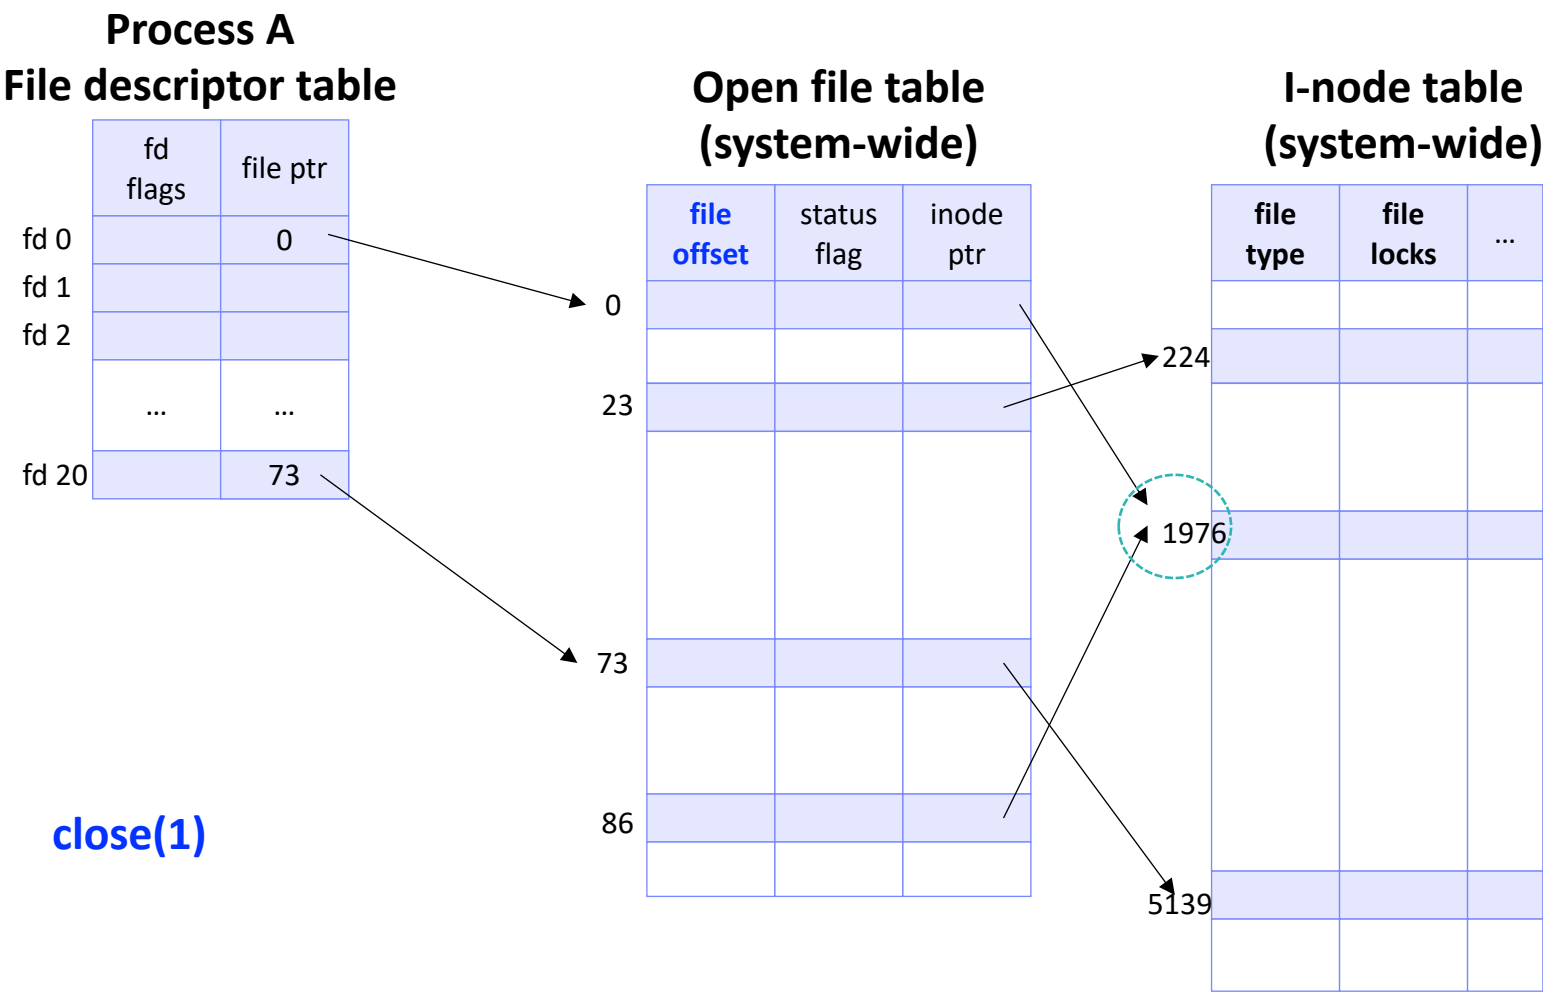
Duplicating File Descriptors



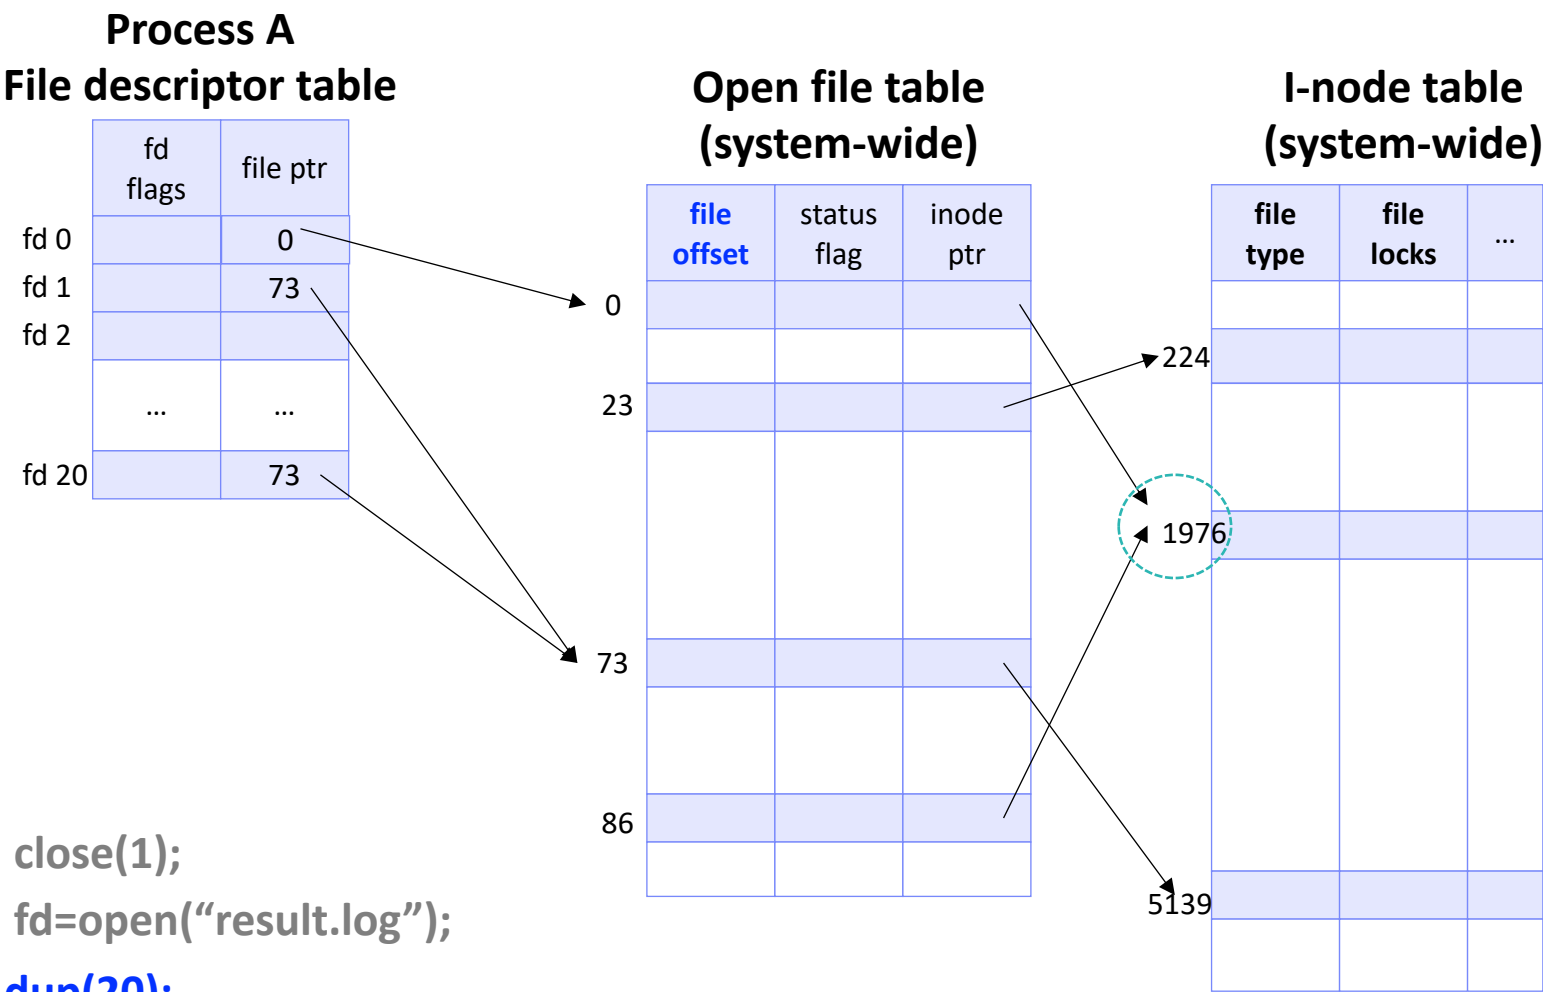
Duplicating File Descriptors



Duplicating File Descriptors



Duplicating File Descriptors



`./myscript > result.log`

Ex4. whotofile.c

```
#include <stdio.h>
#include <unistd.h> /* for execlp */
#include <stdlib.h> /* for exit */

main()
{
    int pid;
    int fd;

    printf("About to run who into a file\n");

    /* create a new process or quit */
    if( (pid = fork() ) == -1 ){
        perror("fork");
        exit(1);
    }

    /* child does the work */
    if( pid == 0 ){
        close(1); /* close */
        fd = creat("userlist", 0644); /* then open */
        execlp("who", "who", NULL); /* and run */
        perror("execlp");
        exit(1);
    }

    /* parent waits then reports */
    if( pid != 0 ){
        wait(NULL);
        printf("Done running who. Results in userlist\n");
    }
}
```

Summary of Redirection to Files

- Three basic facts

- Standard input, output, and error are file descriptors 0, 1, and 2
- The kernel always uses the lowest numbered unused file descriptor
- The set of file descriptors is passed unchanged across exec calls

- The shell also supports the following forms:

- `who > userlog` → write
- `who >> userlog` → add
- `sort < data` → read