

# Ch11. Connecting to Processes Near and FAR - Servers and Sockets

Prof. Seokin Hong

Kyungpook National University

Fall 2019

# Chapter Summary

---

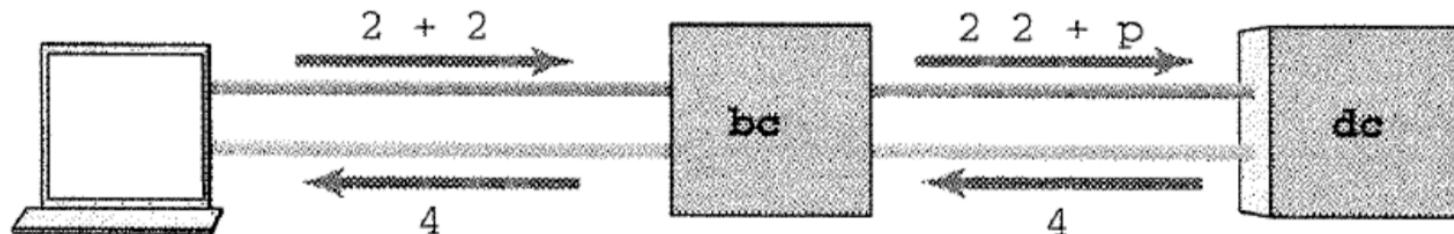
- Client/Server programming using pipes and sockets
- Interprocess communication and client/server design
- The idea and techniques of socket programming

# bc: UNIX calculator

## ■ bc

```
[seokin@compasslab1:~$ bc
bc 1.07.1
Copyright 1991-1994, 1997, 1998, 2000, 2004, 2006, 2008, 2012-2017 Free Software Foundation, Inc.
This is free software with ABSOLUTELY NO WARRANTY.
For details type 'warranty'.
[12^123
54864732218924221509340821667217308113486679281006716245125170218434\
56541709523359082780720277398867836972367369456704108169294512128
[1000+1000
2000
[2000000+200000
20200000
```

## ■ bc is not a calculator



# bc: UNIX calculator

---

## ■ Ideas from bc

- **Client/Server Model**

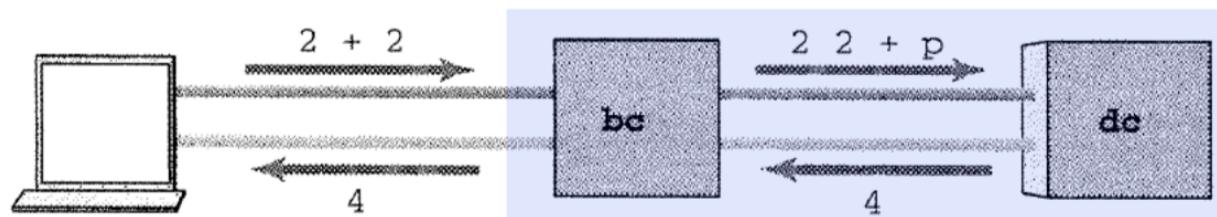
- **dc** provides a service (calculation)
- **bc** provides a user interface and uses the service
- **bc** is called a client of dc

- **Bidirectional Communication**

- Using pipes, you need two pipes

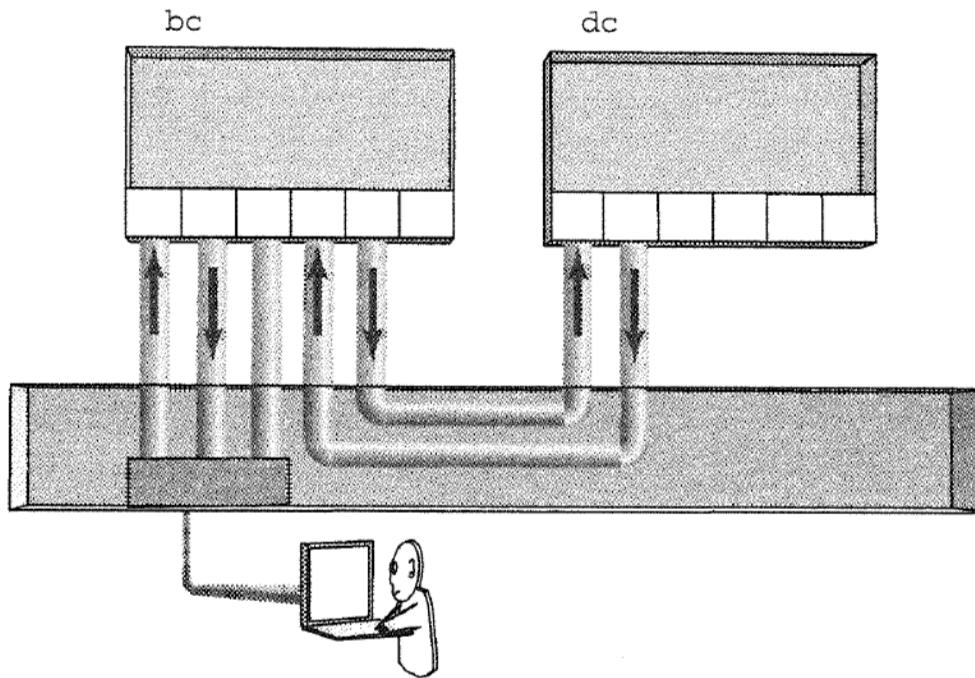
- **Persistent Service (coroutines)**

- Both processes continue to run
- Control passes from one to the other as each completed its part of the jobs
- **bc** has the job of parsing input and printing the computation results
- **dc** has the job of computing



# Coding bc: pipe, fork, dup, exec

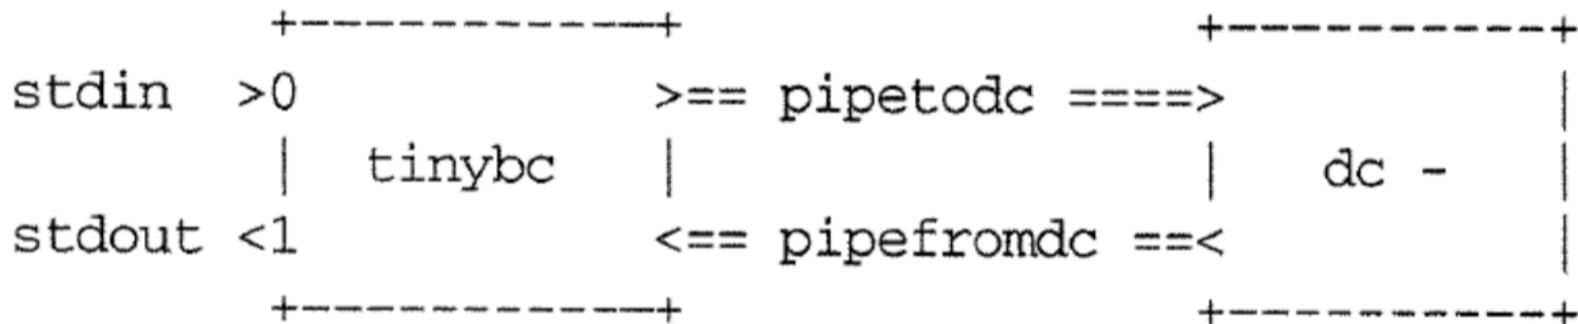
---



- Create two pipes
- Create a process to run **dc**
- In the new process, redirect stdin and stdout to the pipes, then exec **dc**
- In the parent, read and parse user input, write commands to **dc**, read response from **dc**, and send response to user

# Coding bc: pipe, fork, dup, exec

---



## ■ Program outline

- Get two pipes
- fork (get another process)
- In the dc-to-be process,
  - Connect stdin and out to pipes
  - Then exec **dc**
- In the tinybc-process,
  - Talk to human via normal I/O and send stuff via pipe
- Close pipe and dc dies

# Ex1. tinybc.c

---

```
#include    <stdio.h>
#include    <unistd.h>
#include    <stdlib.h>
#include    <fcntl.h>
#include <sys/wait.h>
#define oops(m,x)  { perror(m); exit(x); }

void be_dc(int in[2], int out[2]);
void be_bc(int todc[2], int fromdc[2]);
void fatal(char mess[]);

int main()
{
    int pid, todc[2], fromdc[2];      /*equipment */

    /* make two pipes */
    if( pipe(todc) == -1 || pipe(fromdc) == -1 )
        oops("pipe failed", 1);

    // test
    printf("todc[0] : %d, todc[1] : %d \n", todc[0], todc[1]);
    printf("fromdc[0]: %d, fromdc[1]: %d\n", fromdc[0], fromdc[1]);

    /* get a process for user interface */
    if(( pid = fork()) == -1 )
        oops("cannot fork", 2);
    if( pid == 0 )                  /* child is dc */
        be_dc(todc, fromdc);
    else {
        be_bc(todc, fromdc);       /* parent is ui */
        wait(NULL);               /* wait for child */
    }
}
```

## Ex1. tinybc.c

---

```
void be_dc(int in[2], int out[2])
/*
 *  set up stdin and stdout, then execl dc
 */
{
    /* setuup stdin from pipein */
    if( dup2(in[0], 0) == -1 ) /* copy read end to 0 */
        oops("dc: cannot redirect stdin", 3);
    close(in[0]);           /* moved to fd 0 */
    close(in[1]);           /* won't write here */

    /* setup stdout to pipeout */
    if( dup2(out[1], 1) == -1) /* dupe write end to 1 */
        oops("dc: cannot redirect stdout", 4);
    close(out[1]);          /* moved to fd 1 */
    close(out[0]);          /* won't read from here */

    /* now execl dc with the - option */
    execlp("dc", "dc", "-", NULL);
    oops("Cannot run dc", 5);
}
```

# Ex1. tinybc.c

```
void be_bc( int todc[2], int fromdc[2])
{
    int num1, num2;
    char    operation[BUFSIZ], message[BUFSIZ], *fgets();
    FILE   *fpout, *fpin, *fdopen();

    /* setup */
    close(todc[0]);           /* won't read from pipe to dc */
    close(fromdc[1]);         /* won't write to pipe from dc */

    fpout = fdopen( todc[1], "w" ); /* convert file descriptors to stream */
    fpin  = fdopen( fromdc[0], "r" );

    if( fpout == NULL || fpin == NULL )
        fatal("Error converting pipes to streams");

    /* main loop */
    while( printf("tinybc: "), fgets(message, BUFSIZ, stdin) != NULL )
    {
        /* parse input */
        if( sscanf(message, "%d%*[+*/^]%" , &num1, operation, &num2) != 3)
        {
            printf("syntax error\n");
            continue;
        }

        if( fprintf( fpout, "%d\n%d\n%c\np\n", num1, num2, *operation ) == EOF )
            fatal("Error writing");
        fflush( fpout );

        if( fgets( message, BUFSIZ, fpin ) == NULL )
            break;
        printf("%d %c %d = %s", num1, *operation, num2, message ); // stdout
    }
    fclose(fpout); /* close pipe */
    fclose(fpin);  /* dc will see EOF */
}

void fatal( char mess[] )
{
    fprintf(stderr, "Error: %s\n", mess);
    exit(1);
}
```

## Ex1. tinybc.c

---

```
/* main loop */
while( printf("tinybc: "), fgets(message, BUFSIZ, stdin) != NULL )
{
    /* parse input */
    if( sscanf(message, "%d%[-+*/^] %d", &num1, operation, &num2) != 3)
    {
        printf("syntax error\n");
        continue;
    }

    if( fprintf( fpout, "%d\n%d\n%c\np\n", num1, num2, *operation ) == EOF )
        fatal("Error writing");
    fflush( fpout );

    if( fgets( message, BUFSIZ, fpin ) == NULL )
        break;
    printf("%d %c %d = %s", num1, *operation, num2, message ); // stdout
}
fclose(fpout); /* close pipe      */
fclose(fpin); /* dc will see EOF  */

void fatal( char mess[] )
{
    fprintf(stderr, "Error: %s\n", mess);
    exit(1);
}
```

# **fdopen: Making file descriptors look**

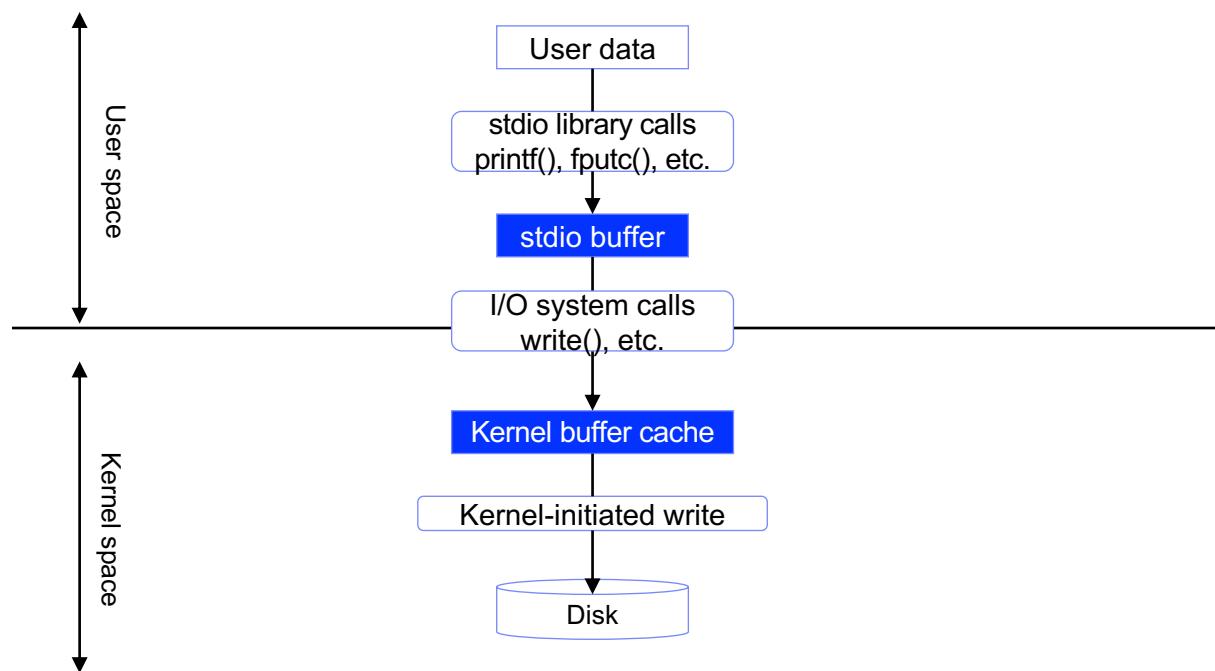
---

## ■ ***File stream***

- A pointer to a FILE structure that contains information about a file.
- permits user-controllable buffering and formatted input and output.

## ■ **fopen: file name → FILE \***

## ■ **fdopen : file descriptor → FILE \***



# Contents

---

- **bc: a Unix Calculator**
- **popen: Making Process Look like Files**
- **Sockets: Connecting to Remote Processes**

# What `popen` does

---

- **`fopen()` opens a buffered connection to a file**

```
FILE *fp;  
fp = fopen("file1", "r");  
c = getc(fp);  
fgets(buf, len, fp);  
fscanf(fp, "%d%d%s", &x, &y, x);
```

- **`popen()` opens a buffered connection to a process**

- Implemented using `fork()`, `pipe()`, `exec()`

```
File *fp;  
fp = popen("ls", "r");  
fgets(buf, len, fp);  
pclose(fp);
```

# What popen does

- Similarities between popen and fopen

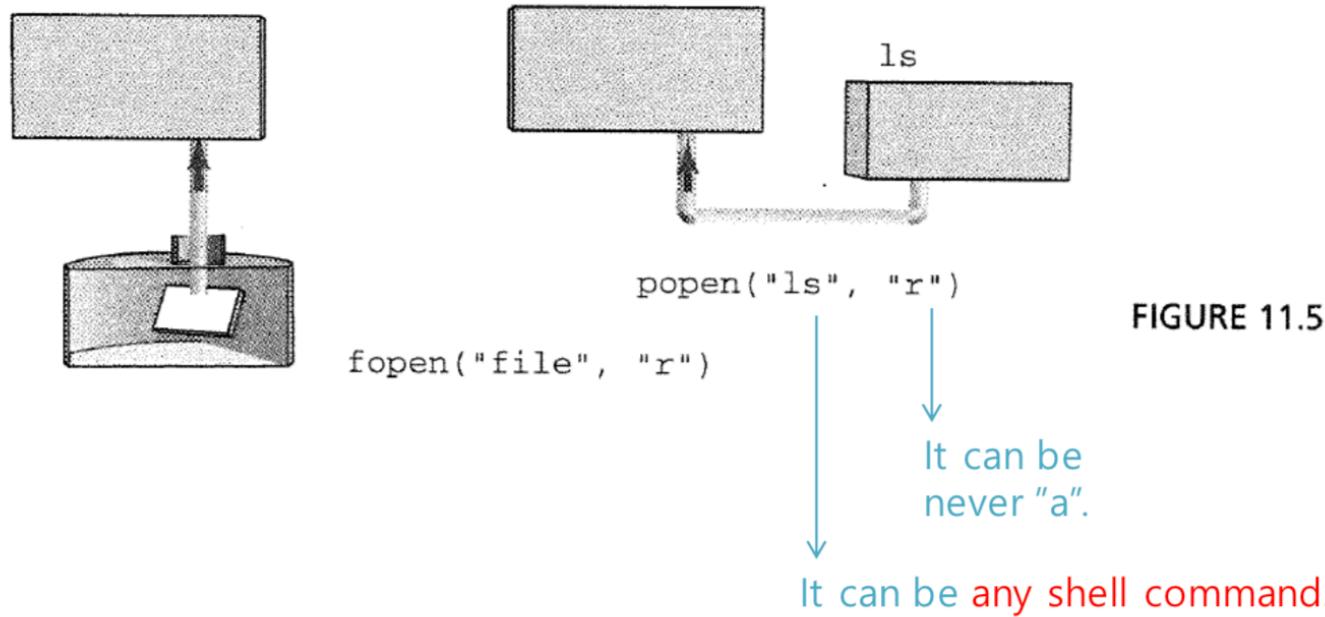


FIGURE 11.5

- **pclose()** is required

- A process needs to wait while the child process executes
- `pclose()` calls `wait()`

# Ex2: popendemo.c

# Writing popen()

## ■ How does popen work?

- popen runs a program and returns a connection to the standard input or standard output of that program.

```
popen(cmd, "r")
```

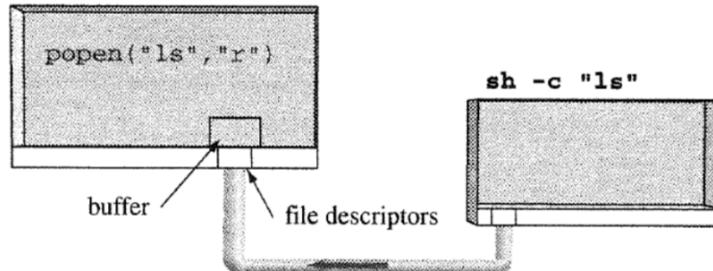
```
pipe(p)  
fork()
```

Parent

Child

```
close(p[1])  
fp = fdopen(p[0],"r");  
return fp;
```

```
close(p[0])  
dup2(p[1],1);  
close(p[1]);  
execl("/bin/sh", "sh", "-c", cmd, NULL);
```



## Ex3. popen.c

---

```
#include    <stdio.h>
#include    <signal.h>
#include    <stdlib.h>
#include    <unistd.h>

#define READ      0
#define WRITE     1

FILE *popen(const char *command, const char *mode)
{
    int pfp[2], pid;          /* the pipe and the process */
    FILE   *fdopen(), *fp;    /* fdopen makes a fd a stream */
    int parent_end, child_end; /* of pipe */

    if (*mode == 'r') {        /* figure out direction */
        parent_end = READ;
        child_end = WRITE ;
    } else if ( *mode == 'w' ){
        parent_end = WRITE;
        child_end = READ ;
    } else return NULL ;

    if ( pipe(pfp) == -1 )           /* get a pipe */
        return NULL;
    if ( (pid = fork()) == -1 ){     /* and a process */
        close(pfp[0]);            /* or dispose of pipe */
        close(pfp[1]);
        return NULL;
    }

    /* ----- parent code here ----- */
    /* need to close one end and fdopen other end */

    if ( pid > 0 ){
        if (close( pfp[child_end] ) == -1 )
            return NULL;
        return fdopen( pfp[parent_end] , mode); /* same mode */
    }
}
```

## Ex3. popen.c

---

```
/* ----- child code here ----- */
/*   need to redirect stdin or stdout then exec the cmd */

if ( close(pfp[parent_end]) == -1 ) /* close the other end */
    exit(1);                      /* do NOT return */

if ( dup2(pfp[child_end], child_end) == -1 )
    exit(1);

if ( close(pfp[child_end]) == -1 ) /* done with this one */
    exit(1);
                                /* all set to run cmd */
execl( "/bin/sh", "sh", "-c", command, NULL );
exit(1);
}

int main()
{
    FILE    *fp;
    char    buf[BUFSIZ];

    fp = popen("ls", "r");
    while( fgets(buf, BUFSIZ, fp) != NULL)
        fputs(buf, stdout);

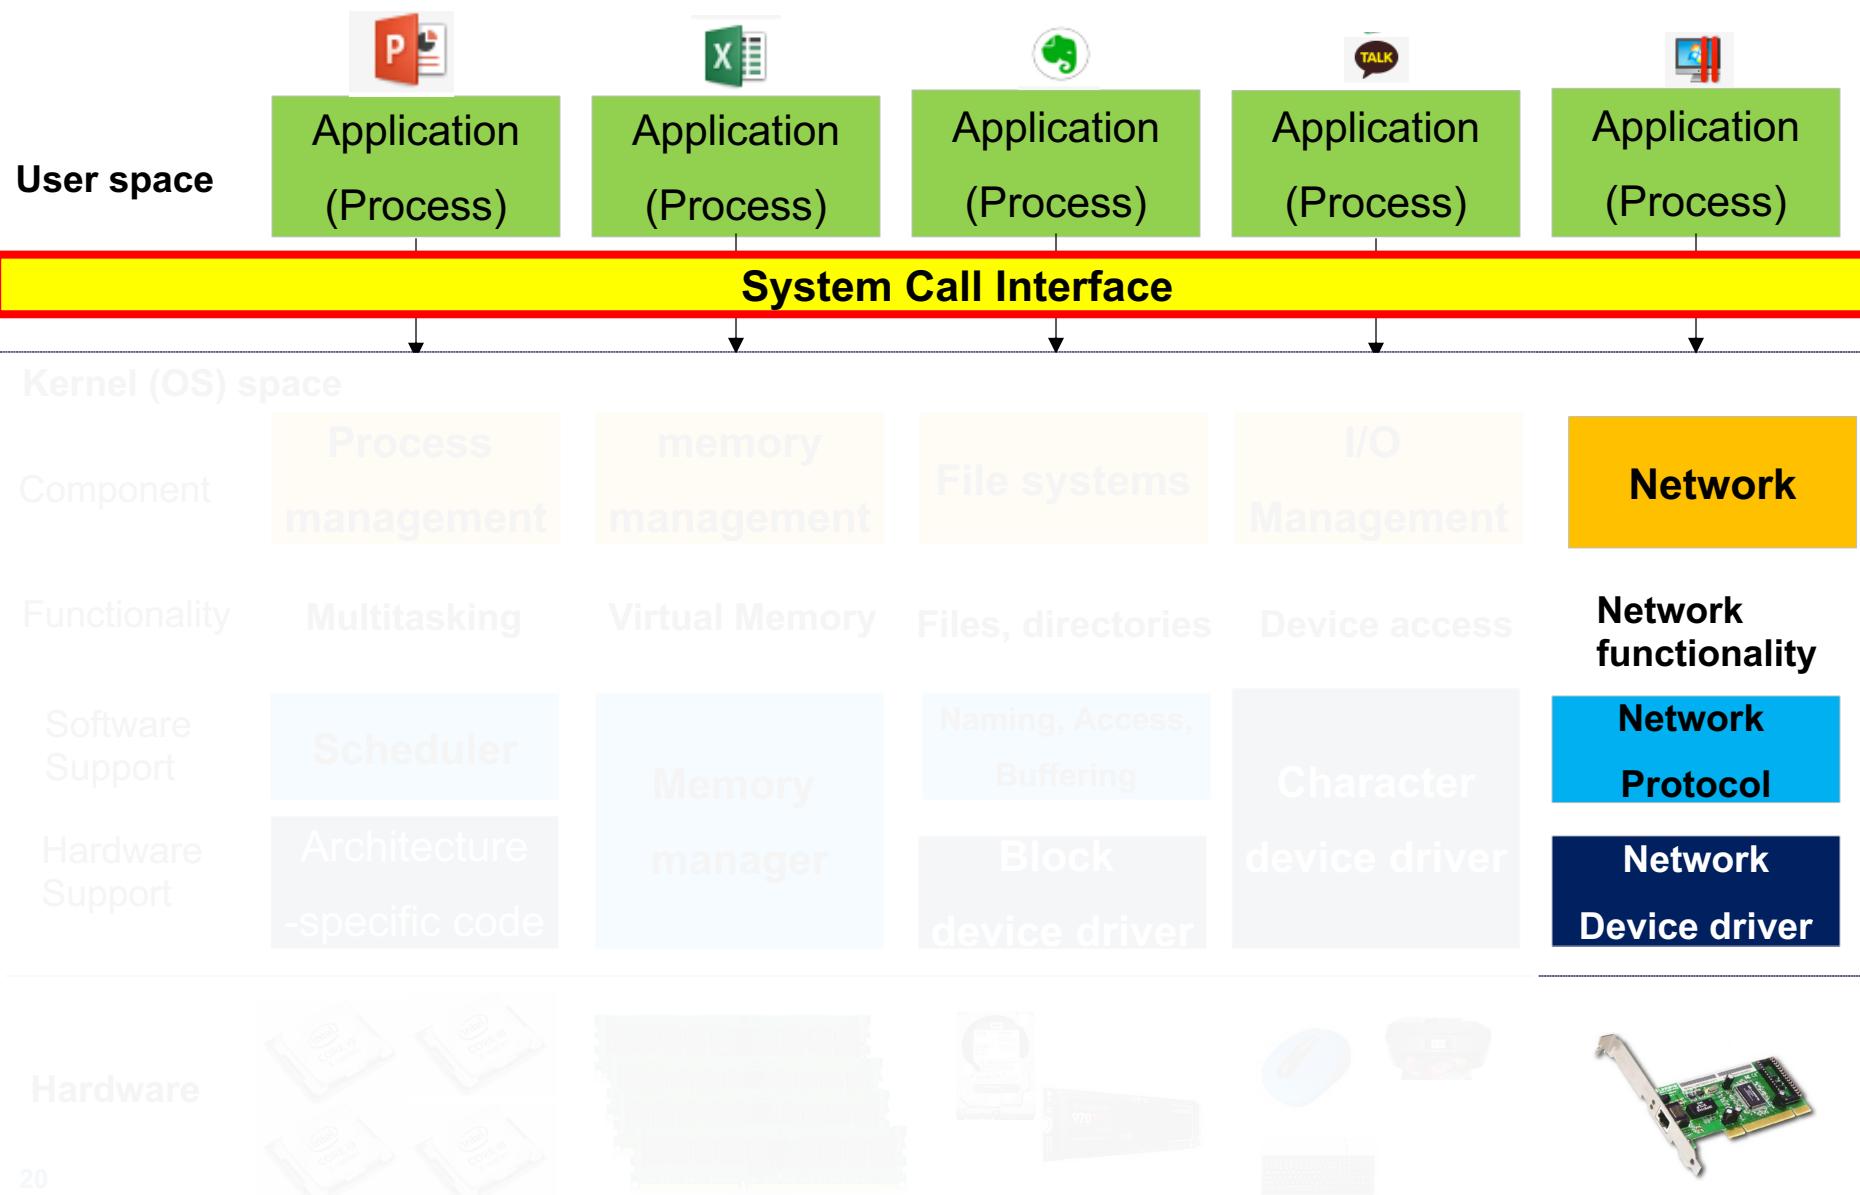
    return 0;
}
```

# Contents

---

- **bc: a Unix Calculator**
- **popen: Making Process Look like Files**
- **Sockets: Connecting to Remote Processes**

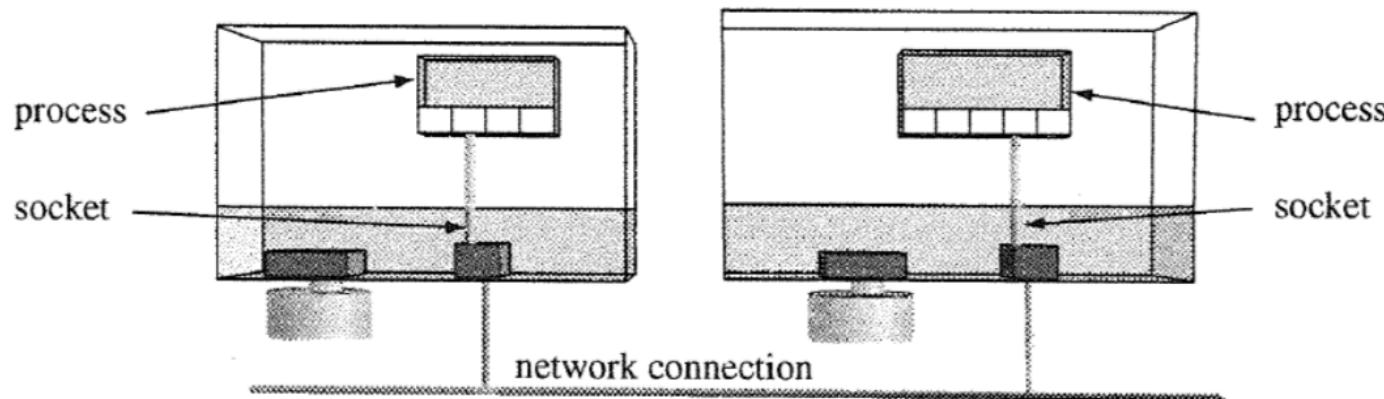
# Sockets: Connecting to Remote Processes



# Sockets

---

- Limitations of pipes
  - It can only connect related processes: ex) parent/child
  - It can only connect processes on the same computer
- Unix/Linux provides another method of interprocess communication: **sockets**
- Sockets allow processes to create pipe-like connections to unrelated processes and even to process on other computers



# Important Concepts in Socket Programming

---

## ■ Client and server

- In Unix terms, server is a program, not a computer
- A server process waits for a request, processes it and loops back to take the next request

## ■ Hostname and port:

- A server on the internet is a process running on a computer
- The computer is called the host, and it has a hostname.
- A port number is allocated to each server program.

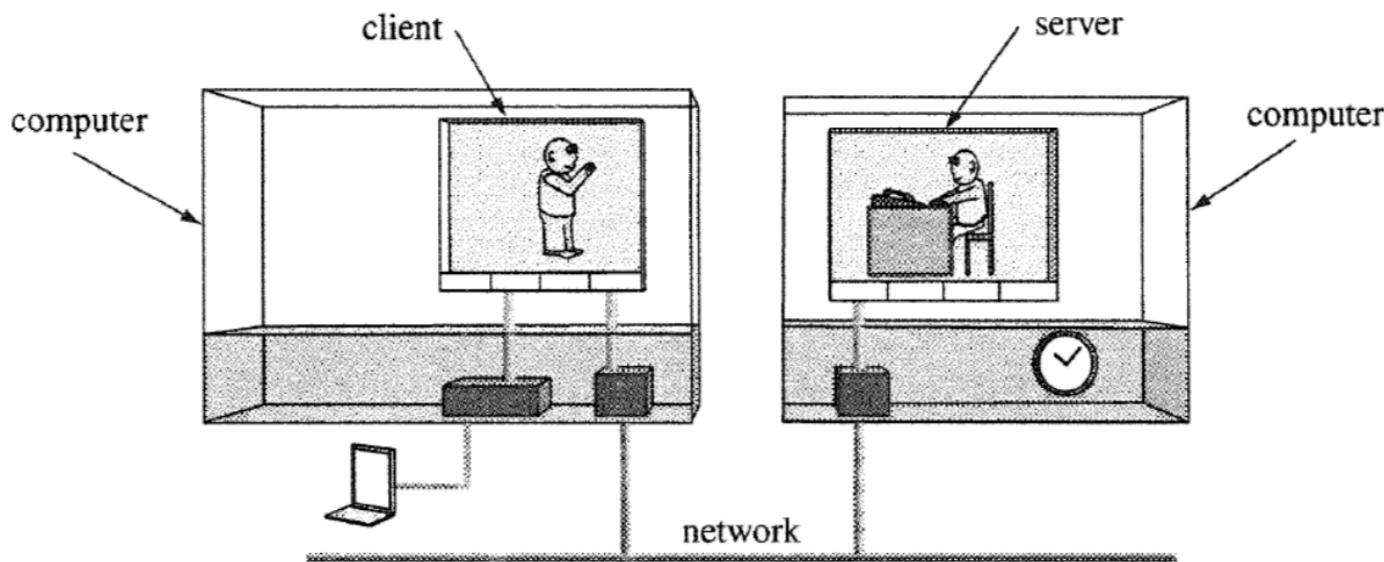
## ■ Address

- Like phone-number or zip code, a host has an address

## ■ Protocol

- A set of rules that governs interaction between client and server

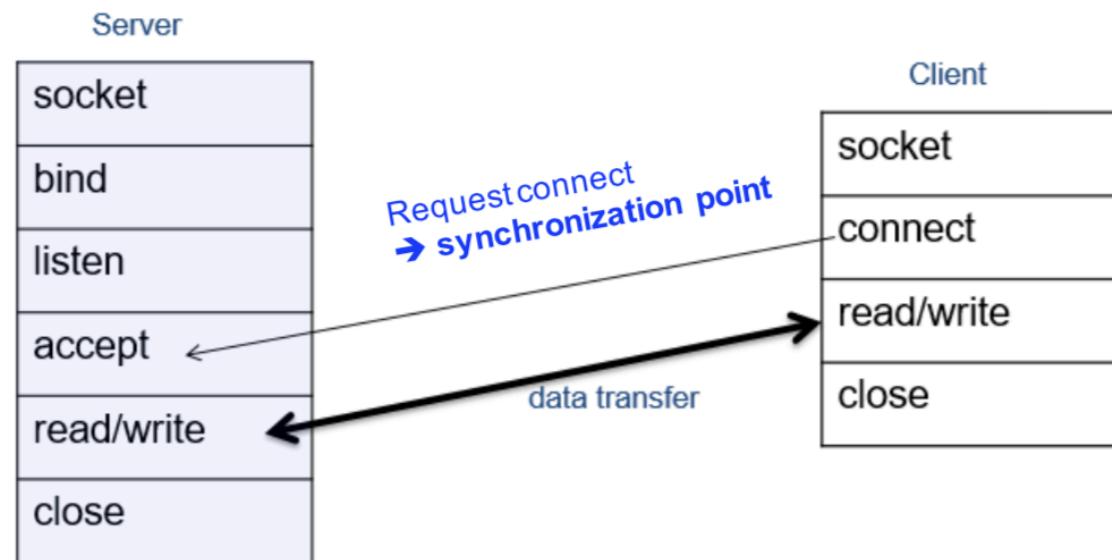
# Example: Time Server and Client



```
[seokin@compasslab2 ch11]$ ./timeserv &
[1] 28592
[seokin@compasslab2 ch11]$ ./timeclnt compasslab2 13000
Wow! got a call!
The time here is ..Mon Nov 18 10:39:30 2019
```

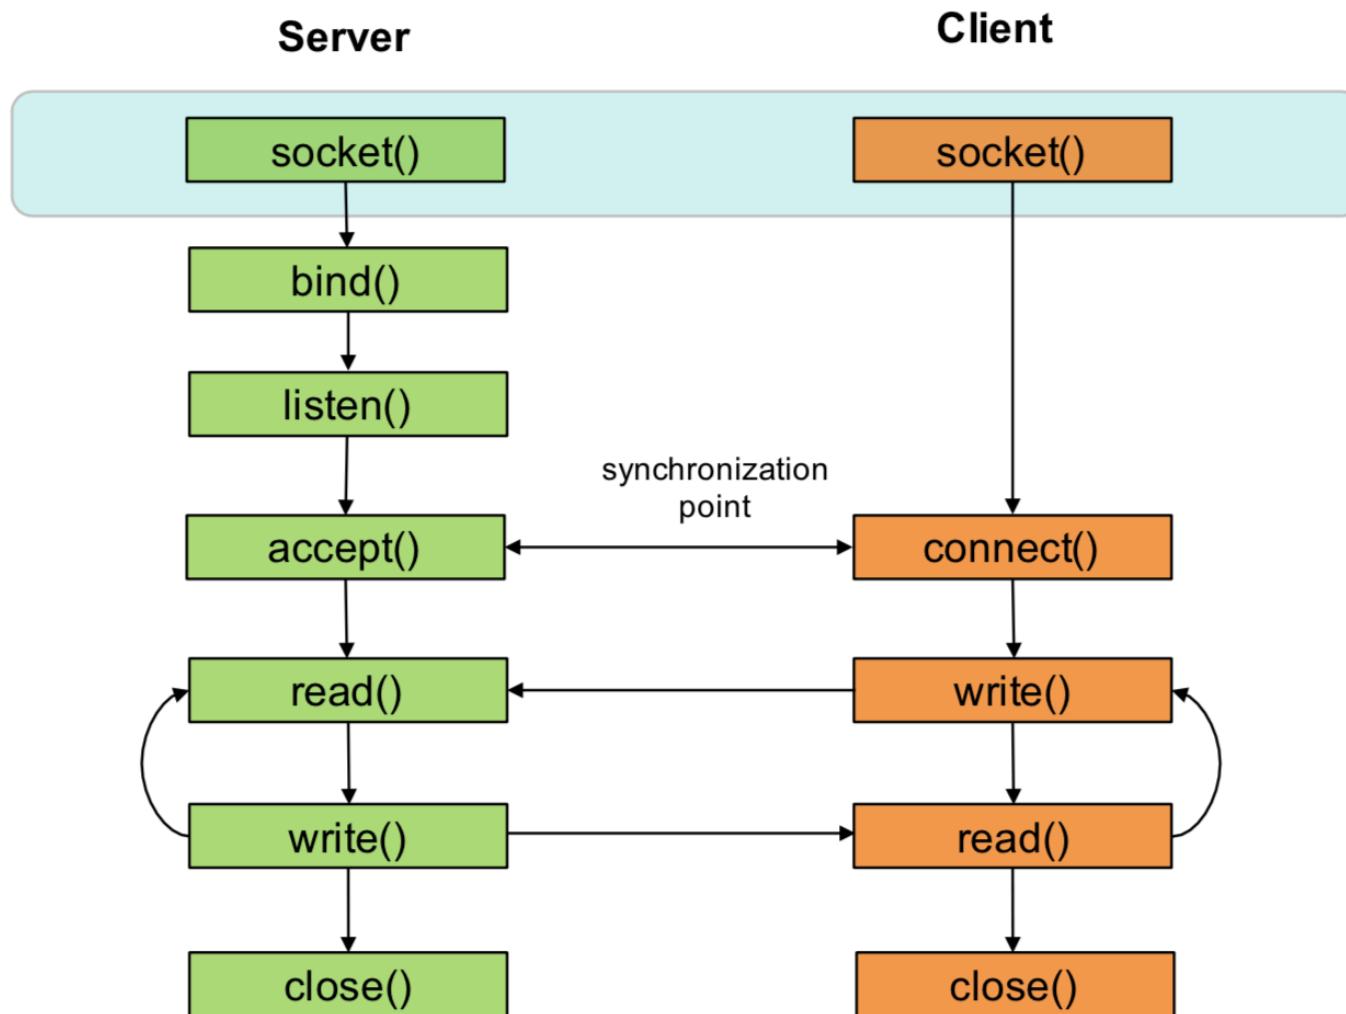
# Writing timeserv.c : A Timer Server

Action	Syscall
Get a phone line	socket()
Assign a number	bind()
Allow incoming calls	listen()
Wait for a call	accept()
Transfer data	read()/write()
Hang up	close()



# Writing timeserv.c : A Timer Server

- Step 1: Create a socket



# Writing timeserv.c : A Timer Server

---

## ■ Step 1: Create a socket → `socket()`

```
#include <sys/socket.h>
```

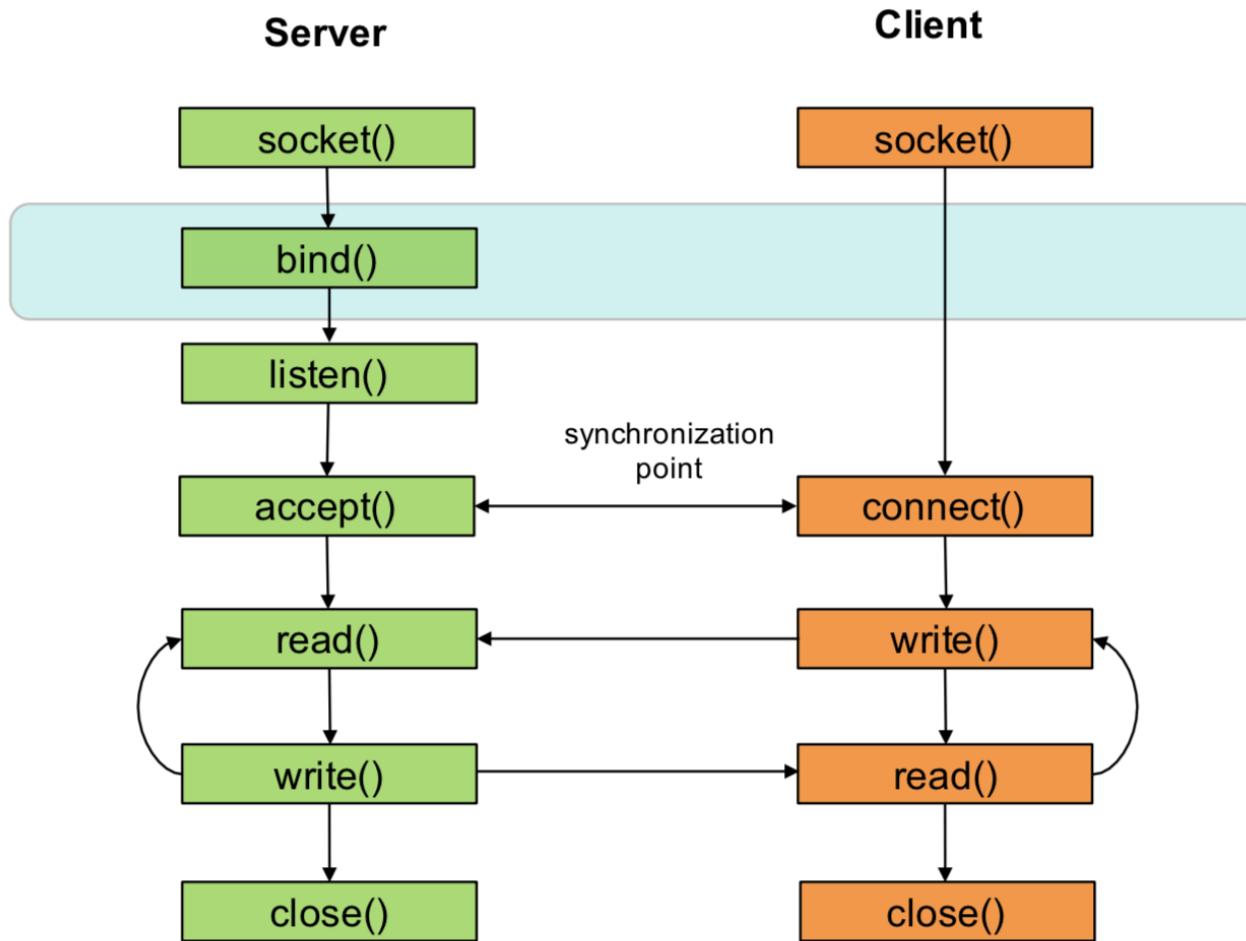
```
int socket(int domain, int type, int protocol);
```

Returns file descriptor on success, or -1 on error

- **domain** : communication domain (PF\_INET or PF\_UNIX)
- **type** : socket type (SOCK\_STREAM or SOCK\_DGRAM)
- **protocol** :
  - Specify 0 (for SOCK\_STREAM and SOCK\_DGRAM)
  - Nonzero for some socket types (e.g., SOCK\_RAW)
- **Return value** : file descriptor

# Writing timeserv.c : A Timer Server

- Step 2: Bind a socket to an address



# Writing timeserv.c : A Timer Server

---

- Step 2: Bind a socket to an address → **bind()**

```
#include <sys/socket.h>
```

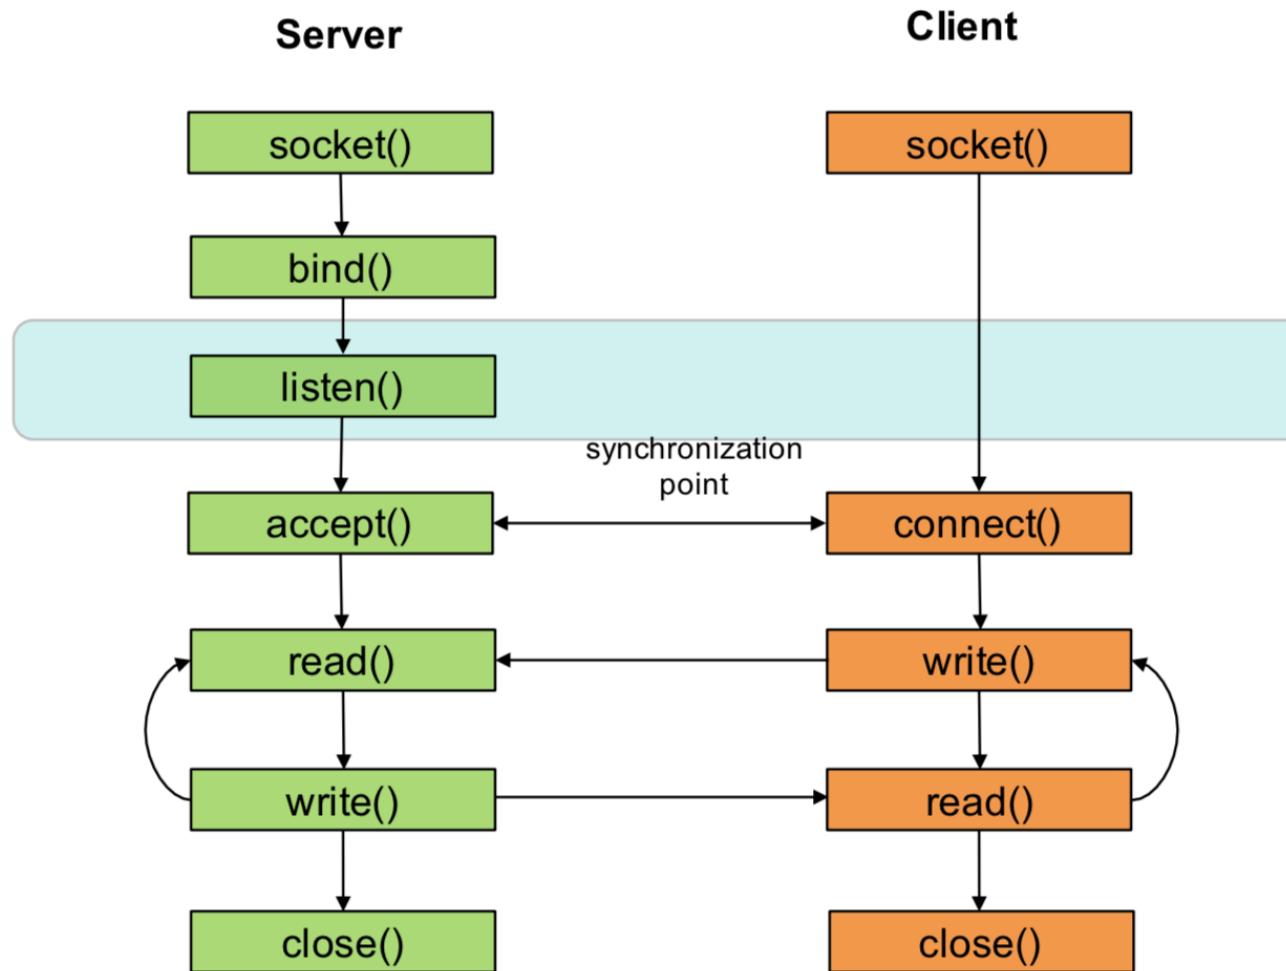
```
int bind(int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

Returns 0 on success, or -1 on error

- **sockfd** : file descriptor obtained from `socket()`
- **addr** : a pointer to a structure specifying the address to which this socket is to be bound
- **addrlen** : the size of the address structure

# Writing timeserv.c : A Timer Server

- Step 3: Allow incoming calls on socket



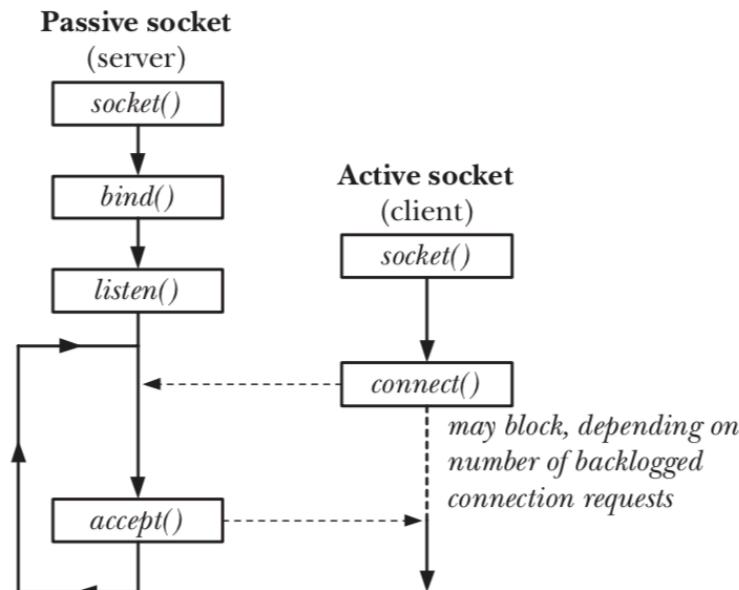
# Writing timeserv.c : A Timer Server

## ■ Step 3: Allow incoming calls on socket → **listen()**

- Marks the socket referred to by the file descriptor as passive

```
#include <sys/socket.h>
int listen(int sockfd, int backlog);           //Returns 0 on success, or -1 on error
```

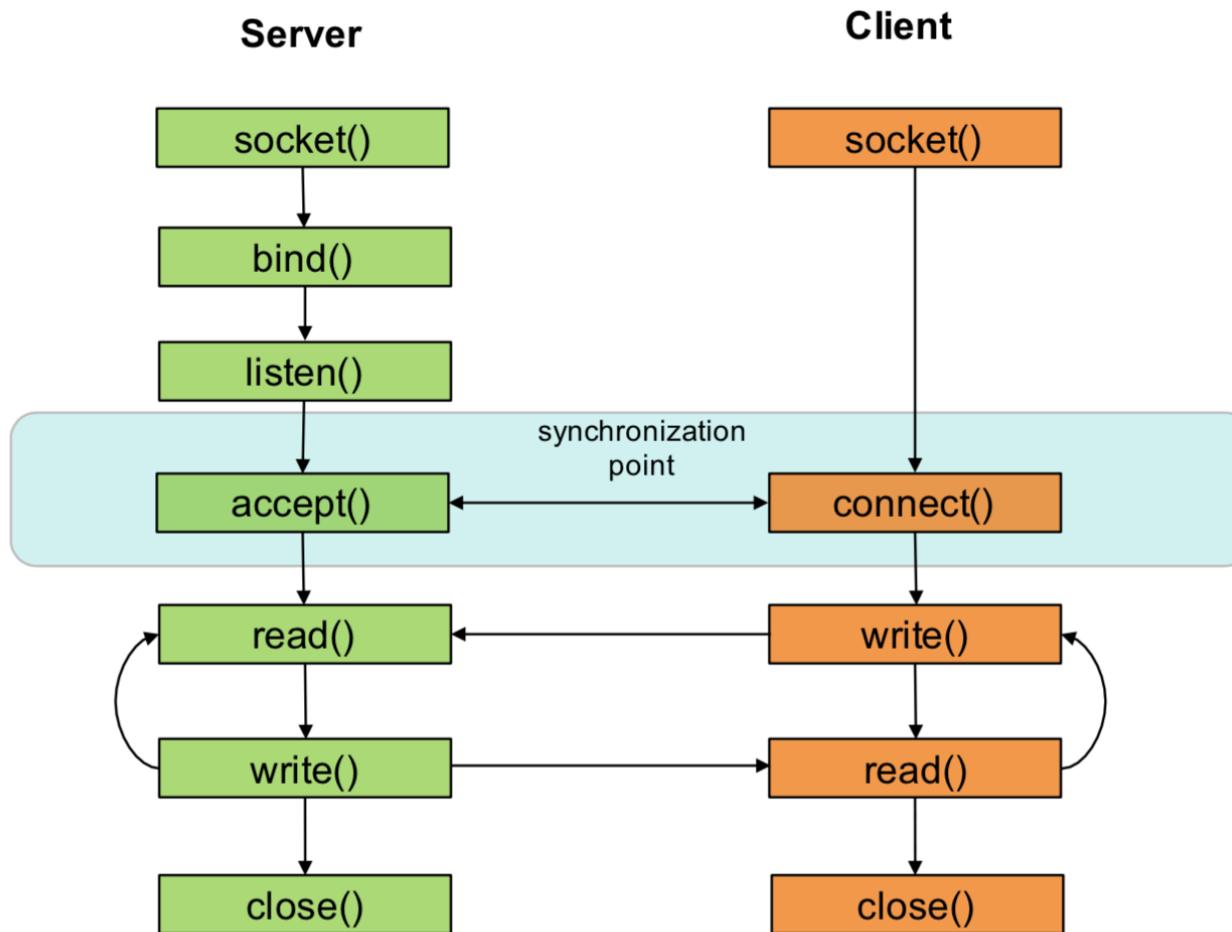
- **sockfd** : file descriptor obtained from `socket()`
- **backlog** : used to limit the number of pending request



Client may call `connect()` before server calls `accept()`  
→ the connection is pended  
→ kernel must record some information about each pending connection

# Writing timeserv.c : A Timer Server

- Step 4: Wait for and accept a connection



# Writing timeserv.c : A Timer Server

---

## ■ Step 4: Wait for and accept a connection → **accept()**

- accept an incoming connection on the listening socket
- If no pending connections, the accept() is **blocked**

```
#include <sys/socket.h>
```

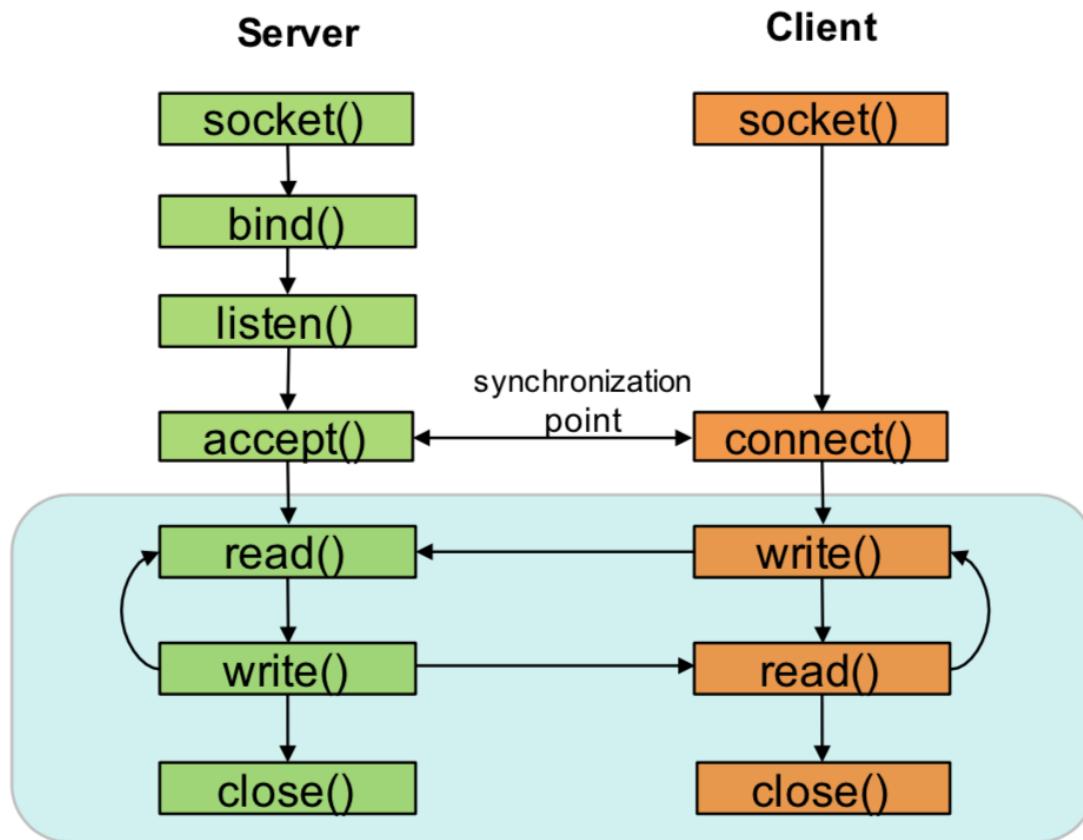
```
int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen);
```

Returns file descriptor on success, or -1 on error

- **sockfd** : file descriptor obtained from socket()
- **addr** : pointer to structure that specify address of the peer address
- **addrlen** : pointer to the size of the address structure
- If we are not interested in the address of the peer socket, addr and addrlen are specified as NULL and 0, respectively.
- **Return value** : **a new socket** that is connected to the peer socket
  - The listening socket remain open and can be used to accept further connection

# Writing timeserv.c : A Timer Server

- Step 5: Transfer data → `read()`, `write()`
- Step 6: Close connection → `close()`



## Ex4: timeserv.c

---

```
/* timeserv.c - a socket-based time of day server
 */
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <netdb.h>
#include <time.h>
#include <string.h>
#include <stdlib.h>

#define PORTNUM 13000 /* our time service phone number */
#define HOSTLEN 256
#define oops(msg) { perror(msg); exit(1); }

int main(int ac, char *av[])
{
    struct sockaddr_in saddr; /* build our address here */
    struct hostent *hp; /* this part of our */
    char hostname[HOSTLEN]; /* address */
    int sock_id, sock_fd; /* line id, file descriptor */
    FILE *sock_fp; /* use socket as steam */
    char *ctime(); /* convert secs to string */
    time_t thetime; /* the time we report */

    /*
     * Step 1: ask kernel for a socket
     */
    sock_id = socket( PF_INET, SOCK_STREAM, 0 );
    if( sock_id == -1)
        oops( "socket" );
```

## Ex4: timeserv.c

---

```
/*
 * Step 2: bind address to socket. Address is host, port
 */
bzero( (void *) &saddr, sizeof( saddr) ); /* clear out struct */

gethostname( hostname, HOSTLEN );      /* where am I ? */
hp = gethostbyname( hostname );       /* get info about host */
                                       /* fill in host part */
bcopy( (void *) hp->h_addr, (void *) &saddr.sin_addr, hp->h_length );
saddr.sin_port = htons(PORTNUM);      /* fill in socket port */
saddr.sin_family = AF_INET;           /* fill in addr family */

if( bind(sock_id, (struct sockaddr *) &saddr, sizeof(saddr)) != 0 )
    oops( "bind" );

/*
 * Step 3: allow incoming calls with Qsize=1 on socket
 */

if ( listen (sock_id, 1) != 0 )
    oops( "listen" );

/*
 * main loop: accept(), write(), close()
 */

while (1) {
    sock_fd = accept (sock_id, NULL, NULL); /* wait for call */
    printf("Wow! got a call!\n");
    if( sock_fd == -1 )
        oops( "accept" ); /* error getting calls */

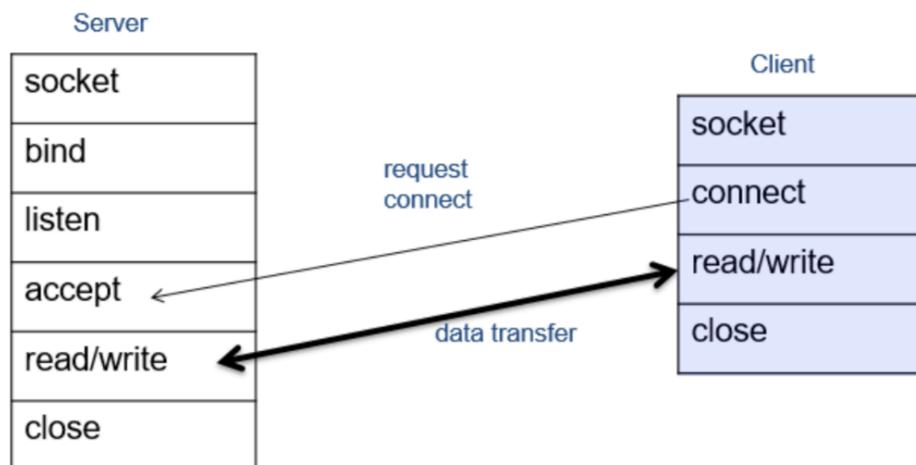
    sock_fp = fdopen( sock_fd, "w");/* we'll write to the */
    if( sock_fp == NULL )          /* socket as a stream */
        oops ("fdopen" ); /* unless we can't */

    thetime = time(NULL);         /* get time */
                                /* and convert to string */
    fprintf( sock_fp, "The time here is .." );
    fprintf( sock_fp, "%s", ctime(&thetime) );
    fclose( sock_fp );           /* release connection */
}
```

# Writing timeInt.c: a Time Client

---

action	syscall
1. Get a phone line	socket()
2. Call the server	connect()
3. Transfer data	read()/write()
4. Hang up	close()



# Writing timeInt.c: a Time Client

---

- Establish a connection : **connect()**

- connect the active socket to the listening socket

```
#include <sys/socket.h>

int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

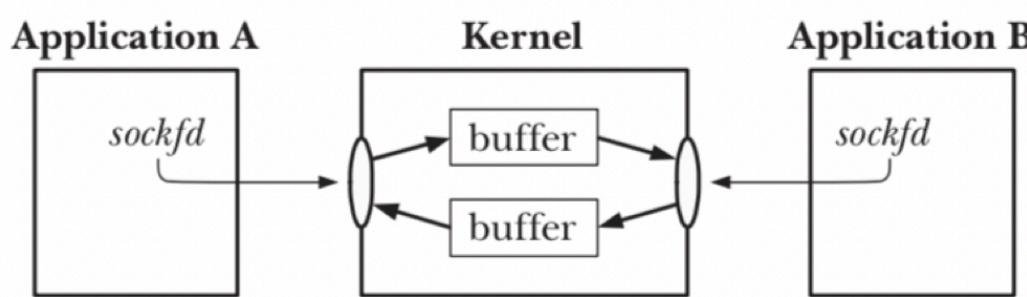
Returns 0 on success, or -1 on error

- **sockfd** : file descriptor of the active socket obtained from socket()
  - **addr** : pointer to structure that specifies address of the listening socket
  - **addrlen** : size of the address structure

# Writing timeInt.c: a Time Client

---

- Use `read()` and `write()` system call to perform I/O
  - Both calls may be used on each end of the connection since the sockets are bidirectional



- After a socket is closed with `close()`
  - `read()` returns zero (the end-of-file)
  - `write()` fails with the error EPIPE and receives a SIGPIPE signal

## Ex5: timeclnt.c

---

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <netdb.h>

#include <stdlib.h>
#include <strings.h>

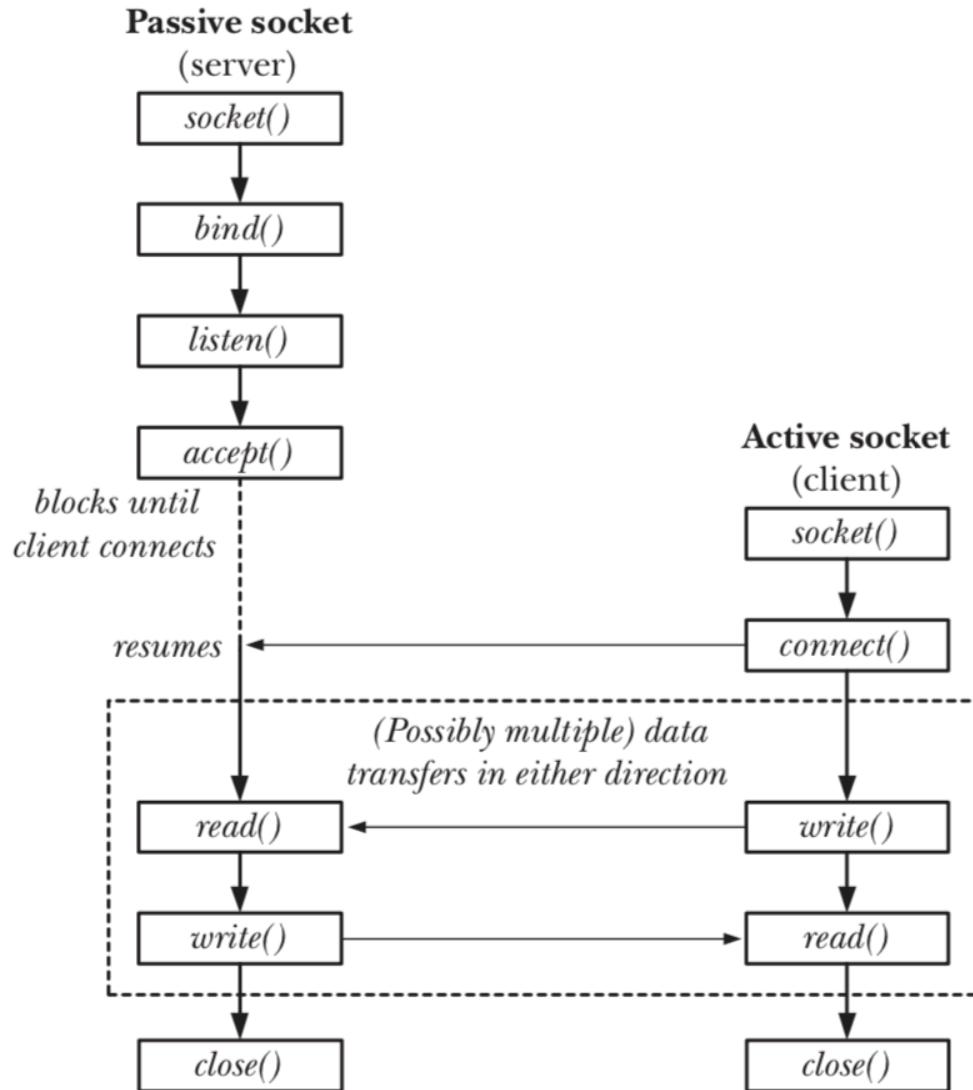
#define oops(msg) { perror(msg); exit(1); }

int main(int ac, char *av[])
{
    struct sockaddr_in servaddr; /* the number to call */
    struct hostent *hp;          /* used to get number */
    int sock_id/*, sock_fd*/;   /* the socket and fd */
    char message[BUFSIZ];       /* to receive message */
    int messlen;                /* for message length */

    /*
     * Step 1: Get a socket
     */
    sock_id = socket( AF_INET, SOCK_STREAM, 0); /* get a line */
    if( sock_id == -1)
        oops( "socket" );                      /* or fail */
```



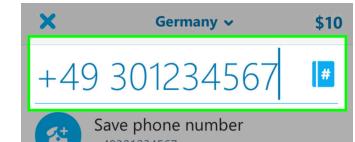
# Summary of system calls used with Socket



`socket()` : create a socket =



`bind()` : bind a socket to address=



`listen()` : notify the kernel of its willingness to accept incoming connections =



`accept()` : notify the kernel of its willingness to accept incoming connections =



`connect()` : establish a connection



# Appendix

---

```
struct hostent {
    char *h_name; /* official name of host */
    char **h_aliases; /* alias list */
    int h_addrtype; /* host address type */
    int h_length; /* length of address */
    char **h_addr_list; /* list of addresses */
};

#define h_addr h_addr_list[0]/* for backward compatibility*/
```

```
struct sockaddr {
    unsigned short sa_family; /* Address family (e.g. AF_INET) */
    char sa_data[14]; /* Family-specific address information */
}
```

```
struct in_addr {
    unsigned long s_addr; /* Internet address (32 bits) */
}
```

```
struct sockaddr_in {
    unsigned short sin_family; /* Internet protocol (AF_INET) */
    unsigned short sin_port; /* Address port (16 bits) */
    struct in_addr sin_addr; /* Internet address (32 bits) */
    char sin_zero[8]; /* Not used */
}
```